

Online Constrained Control of Storage Systems via Simplex Disturbance-Action Policies

Kamiar Asgari and Michael J. Neely

Abstract—We study online control of a scalar storage system with nonnegative adversarial resource arrivals, convex state-action costs, and state-dependent action constraints. At each time, the controller selects a feasible action before the current resource arrival and cost are revealed. To handle the resulting coupling between feasibility and online learning, we introduce *Simplex Disturbance-Action Control* (SDAC), an H -dimensional finite-memory policy class parameterized by a subprobability simplex. The key feature of SDAC is that every fixed policy in the class is feasible by construction. Our adaptive controller updates these parameters by entropic online mirror descent while preserving feasibility along the time-varying trajectory. We prove regret $O(\sqrt{T \log(H+1)})$ relative to the best fixed SDAC policy. As the retention coefficient α approaches one, the leading factor scales as $\Theta((1-\alpha)^{-3/2})$. We also derive a minimax lower bound which, when T is at least on the order of $(1-\alpha)^{-1}$, scales as $\Omega((1-\alpha)^{-1}\sqrt{T})$ near $\alpha = 1$, leaving a factor $(1-\alpha)^{-1/2}$ between the upper and lower bounds. We also show that SDAC geometrically approximates a decreasing-allocation policy class that spreads each arrival over time according to a common decreasing allocation sequence, a structure that arises in stochastic energy-harvesting control. Geometric allocation sequences induce all feasible proportional state-feedback policies. With $H = \Theta(\log T)$, the resulting regret against the decreasing-allocation class, and hence against feasible proportional state-feedback policies, is $O(\sqrt{T \log \log T})$. The framework applies to energy-harvesting batteries and other storage-constrained resource systems.

Index Terms—Online control, disturbance-action policies, storage systems, energy harvesting, regret minimization, state-dependent action constraints.

I. INTRODUCTION

A. Motivation and Problem Setting

MANY dynamic resource systems can be modeled by a storage state that evolves over time and constrains the control action at every time.

This creates a constrained discrete-time control problem with state-action coupling: aggressive early actions can harm later performance, while overly conservative actions can increase the stage cost. Our objective is to remain feasible at every time while competing with the best fixed policy in a tractable class that approximates a broad family of state-responsive allocation rules.

B. Model and Information Structure

We study a discrete-time storage system with state $x_t \geq 0$, where x_t models the storage level at the beginning of time

t . At each time t , the controller chooses a control action u_t under the state-dependent action constraint

$$0 \leq u_t \leq \alpha x_t, \quad (1)$$

where $\alpha \in (0, 1)$ is a fixed known retention coefficient (so a fraction $1 - \alpha$ of the state is lost each period). After the action is chosen, a nonnegative resource arrival w_t is revealed, and the state evolves according to

$$x_{t+1} = \alpha x_t - u_t + w_t.$$

Although this quantity is commonly called a disturbance in the online-control literature, we call $w_t \in [0, 1]$ an *arrival* to emphasize its nonnegativity; we retain the standard names Disturbance-Action Control (DAC) and Simplex Disturbance-Action Control (SDAC). Nonnegativity is essential to the storage interpretation and to the feasibility properties of SDAC.

Unlike stochastic models, our formulation is adversarial and non-anticipatory. At time t , the environment may select w_t and c_t using the history through time $t-1$, but not the current action; the controller then chooses u_t , after which w_t and the entire function c_t are revealed. Both the online and benchmark trajectories must remain feasible.

C. Application Domains

This abstraction captures a broad set of domains; in each case, the action is limited by (1):

- **Energy harvesting and battery management:** x_t is battery state of charge (retention coefficient α after self-discharge), w_t is harvested energy, and u_t is discharge or consumption.
- **Queueing networks and buffer management:** x_t is queue/buffer occupancy, w_t is packet arrivals, and u_t is service.
- **Inventory control:** x_t is on-hand inventory (retention coefficient α after spoilage/deterioration), w_t is supply or replenishment, and u_t is sales.
- **Water-reservoir management:** x_t is stored water (retention coefficient α after evaporation/seepage), w_t is inflow, and u_t is controlled release.

We take energy harvesting as the running interpretation.

D. Contributions

Our contributions are as follows:

- **Policy class.** We introduce SDAC, a DAC-inspired policy class. We prove that every fixed SDAC policy is feasible, its induced loss is convex and Lipschitz in the

K. Asgari and M. J. Neely are with the Department of Electrical and Computer Engineering, University of Southern California, Los Angeles, CA 90089 USA (e-mail: kamiaras@usc.edu; mjneely@usc.edu).

Corresponding author: K. Asgari.

policy parameter, and the class geometrically approximates a *decreasing-allocation policy class*, defined in Section IV-C. Its geometric sequences induce all feasible proportional state-feedback policies.

- **Online control and regret.** We develop an entropic online mirror descent (OMD) controller that updates SDAC parameters over the subprobability simplex while maintaining feasibility of the online trajectory. Our analysis yields $O(\sqrt{T \log(H+1)})$ regret against the best fixed SDAC policy, as stated in Theorem 1. The challenging regime is $\alpha \uparrow 1$, where stored resources persist longer and the leading factor scales as $\Theta((1-\alpha)^{-3/2})$. The SDAC approximation guarantees extend the regret result: choosing $H = \Theta(\log T)$ gives $O(\sqrt{T \log \log T})$ regret against the best decreasing-allocation policy and, as a special case, the best proportional state-feedback policy; see Corollaries 2 and 3.
- **Minimax lower bound.** We prove that $\sqrt{T}$ -order regret is unavoidable for every causal feasible controller, even when the adversarial sequence is deterministic and oblivious; see Theorem 2. For sufficiently long horizons and $\alpha \uparrow 1$, Corollary 1 gives the lower dependence $\Omega((1-\alpha)^{-1})$, whereas Theorem 1 gives the upper dependence $O((1-\alpha)^{-3/2})$. Closing the remaining factor $(1-\alpha)^{-1/2}$ between these dependences is left for future work.

E. Related Work

a) *Online nonstochastic control:* A recent line of work studies online control of linear dynamical systems under adversarial disturbances and costs. Its performance criterion is cumulative regret relative to a benchmark policy class chosen in hindsight, under suitable stability conditions [1], [2]; see [3] for an overview. For known systems in the fully adversarial setting, Simchowit et al. establish the optimal $\sqrt{T}$ horizon dependence without resolving all stability-parameter dependence [4].

A central idea in this literature is DAC, a policy parameterization that combines a stabilizing state-feedback term with a finite linear filter of past disturbances. Under standard assumptions, this representation yields tractable surrogate losses in the policy parameters and supports efficient no-regret updates. Extensions include partial observability [5] and bandit feedback [6].

Constrained online-control methods have also been developed for affine and general static state-input constraints [7], [8]. Our focus is the storage-specific constraint (1), whose feasible set depends on the current state; SDAC handles this dependence with a simplex parameterization.

Beyond algorithm design, our lower-bound analysis studies the joint dependence of regret on the horizon T and the coefficient α . Kumar et al. establish matching memory-dependent upper and lower bounds for online convex optimization with unbounded memory [9]. For discounted memory, their lower bound is of order $\sqrt{T}/(1-\rho)$, paralleling the $\Omega(\sqrt{T}/(1-\alpha))$ rate when the horizon is sufficiently long; see Corollary 1. This agreement is at the level of scaling, while our lower bound

is established directly for state-action costs under the state-dependent storage constraint.

b) *Energy harvesting:* Prior work on online energy harvesting often studies simple rules that adjust power use to the available battery level. Shaviv and Özgür [10] proposed using a fixed fraction of the available energy at each time and proved performance guarantees that hold for any independent and identically distributed (i.i.d.) energy-arrival model. Arafa et al. [11] allowed more general rewards and characterized how energy should be spent after an occasional full recharge; the optimal amounts need not decrease by the same proportion each period. Zibaeenejad and Chen [12] derive a related decreasing spending rule when upcoming arrivals are known in advance. Unlike our setting, these models reveal the current harvested energy before transmission, assume random arrivals from a fixed model, and optimize a fixed reward that depends only on the allocated power, rather than adversarially varying costs that depend jointly on storage state and action.

To demonstrate the expressive power of SDAC, we introduce a decreasing-allocation policy class as a benchmark inspired by this literature. Each policy allocates every arrival over future periods according to the same decreasing sequence, beginning one period later to respect our information pattern. Section IV-C shows that finite-memory SDAC policies approximate every policy in this class with error that decreases geometrically in the memory horizon.

Lin et al. [13] study online age-of-information optimization in an energy-harvesting system with related battery dynamics. They compare with an unrestricted offline optimum through competitive ratio; their objective depends on the energy allocation but not directly on the battery state, and the current input is revealed before allocation. In our model, the controller acts before observing the current arrival and cost, the cost may depend jointly on state and action, and performance is measured by regret against a fixed policy class.

Yu and Neely [14] study expected utility maximization under i.i.d. system states. Asgari and Neely [15] express the guarantee as an $O(\sqrt{T})$ regret-battery-capacity tradeoff and improve its dependence on the decision dimension, allowing arbitrary loss sequences while retaining i.i.d. energy arrivals. Both works use action-only objectives and fixed-decision comparators. Here, arrivals are adversarial, costs depend on both state and action, and the comparator is a fixed state-responsive policy.

c) *Queueing and network optimization:* Queueing control provides a mathematically related perspective based on stochastic arrivals and real-time service decisions. Foundational work emphasizes throughput and queue stability [16]–[18], while stochastic network optimization also studies tradeoffs between time-average objectives and backlog [19]. Our setting differs in its adversarial arrivals, convex state-action costs, and regret benchmark against a policy chosen in hindsight.

In many queueing applications, the natural retention coefficient is $\alpha = 1$ (no abandonment or leakage), so those systems are not direct instances of the $\alpha \in (0, 1)$ model analyzed here. Section VII discusses the additional issues that arise when $\alpha \geq 1$.

Among other results, Khojastepour and Sabharwal [20] study a time-invariant robust scheduler representable as a linear filter of past arrivals. It is designed to minimize average transmit power subject to a maximum packet-delay guarantee, rather than updated online to minimize revealed general costs.

d) Inventory management: Yang et al. [21] study online procurement under inventory constraints. The controller observes the current demand and price, then procures supply while ensuring that demand is met; performance is measured by competitive ratio for aggregate procurement cost. Our controller instead chooses an outflow before observing the current resource arrival and cost, and permits convex state-action costs.

Hihat and Fermanian [22] consider an inventory-control formulation with potentially nonlinear and nondifferentiable state-transition functions and study base-stock policies. While that model is more general in scope, its decision structure differs from ours: as in [21], the controller chooses supply and then observes demand through the transition, whereas our controller chooses an outflow and then observes the resource arrival. Our scalar linear recursion and simplex policy class yield explicit feasibility and regret guarantees; the two frameworks target different dynamics, decisions, and policy classes.

e) Water-reservoir management: Classical dam theory studies storage dynamics with random inflows and controlled releases (see, for example, [23], [24]). Relative to our setting, this literature assumes stochastic inputs and emphasizes stationary or long-run performance quantities. It provides historical context for storage control under state constraints.

F. Paper Organization

Section III formalizes the storage dynamics, information pattern, cost assumptions, feasible domain, and regret benchmark. Section IV then introduces the SDAC policy class. It derives an affine representation of the induced state-action trajectory, uses this representation to establish feasibility and convexity of the surrogate losses, and studies the finite-memory approximation of decreasing-allocation policies, including proportional state feedback as a special case. Section V presents the feasible online controller: a clipped SDAC action coupled with entropic OMD over the subprobability simplex. Section VI proves the regret upper bound, establishes a minimax lower bound, compares their dependence on the retention coefficient, and extends the guarantee to the larger benchmark classes. Section VII outlines extensions to non-contractive retention, vector allocations, and multiple storage states, and Section VIII summarizes the main conclusions. The Appendix contains the deferred approximation, lower-bound, OMD, and stability proofs.

II. NOTATION

For every positive integer n , we write $[n] \triangleq \{1, \dots, n\}$ and $[z]_+ \triangleq \max\{z, 0\}$. For $\mathbf{v} \in \mathbb{R}^n$, the i -th coordinate is v_i ; in particular, the coordinates of $\mathbf{m}_t$ and $\mathbf{g}_t$ are $m_{t,i}$ and $g_{t,i}$. We write $\mathbf{1}$ for the all-ones vector (dimension is clear from context). The Euclidean inner product is denoted by $\langle \cdot, \cdot \rangle$. We

write $\log$ for the natural logarithm. We set $w_\tau \equiv 0$ for $\tau \leq 0$. A fixed SDAC policy with parameter $\mathbf{m}$ is written $\pi_{\mathbf{m}}$, with induced trajectories $x_t(\mathbf{m})$ and $u_t(\mathbf{m})$. A feasible proportional state-feedback policy with gain k induces $x_t(k)$ and $u_t(k)$. A decreasing allocation sequence $\mathbf{q}$ induces $x_t(\mathbf{q})$ and $u_t(\mathbf{q})$. After the online algorithm is introduced, unadorned x_t and u_t denote the realized online trajectory.

III. MODEL AND PROBLEM SETUP

Let T be a positive integer. The system state $x_t \in \mathbb{R}_{\geq 0}$ denotes the resource level at the beginning of time t . The initial state is given as

$$x_1 = 0. \quad (2)$$

Fix a retention coefficient $\alpha \in (0, 1)$. At each time $t \in [T]$, the controller chooses a control action $u_t \in \mathbb{R}_{\geq 0}$ subject to the state-dependent action constraint (1). The system then receives an arrival $w_t \in [0, 1]$ and evolves according to

$$x_{t+1} = \alpha x_t - u_t + w_t. \quad (3)$$

For every feasible trajectory,

$$0 \leq x_{t+1} \leq \alpha x_t + 1.$$

Since $x_1 = 0$, induction gives

$$0 \leq x_t \leq \frac{1 - \alpha^{t-1}}{1 - \alpha} \leq \frac{1}{1 - \alpha}, \quad t \in [T + 1].$$

Thus, the storage level is uniformly bounded over time. In particular, a storage device with capacity $1/(1 - \alpha)$ suffices to realize every feasible trajectory, including the trajectory generated by our online controller, without overflow. Together with the state bound above, (1) implies that every feasible state-action pair belongs to the compact set

$$\mathcal{D}_\alpha \triangleq \left\{ (x, u) : 0 \leq x \leq \frac{1}{1 - \alpha}, 0 \leq u \leq \alpha x \right\}. \quad (4)$$

After the control action is selected, a convex cost function

$$c_t : \mathbb{R}_{\geq 0}^2 \rightarrow \mathbb{R}$$

is revealed, and the incurred cost is $c_t(x_t, u_t)$. We assume each c_t is convex in (x, u) and satisfies the following Lipschitz condition on $\mathcal{D}_\alpha$: there exist constants $L_x, L_u > 0$, independent of t , such that for all $(x, u), (x', u') \in \mathcal{D}_\alpha$,

$$|c_t(x, u) - c_t(x', u')| \leq L_x |x - x'| + L_u |u - u'|. \quad (5)$$

We also assume access, at every $(x, u) \in \mathcal{D}_\alpha$, to a cost subgradient

$$\zeta_t(x, u) = (\zeta_t^x(x, u), \zeta_t^u(x, u)) \in \partial c_t(x, u)$$

satisfying

$$|\zeta_t^x(x, u)| \leq L_x, \quad |\zeta_t^u(x, u)| \leq L_u.$$

The stated subgradient bounds hold, for example, when (5) extends to an open neighborhood of $\mathcal{D}_\alpha$; see [25, Thm. 9.13].

a) *Assumptions:*

- 1) **Retention coefficient:** $\alpha \in (0, 1)$ is fixed and known (a fraction $1 - \alpha$ of the state is lost each period).¹
- 2) **Arrival bounds:** $w_t \in [0, 1]$ for all $t \in [T]$. In particular, resource arrivals are nonnegative, in contrast to signed process-noise disturbances common in classical control.
- 3) **Information pattern:** (w_t, c_t) may depend on the history through time $t - 1$, but not on the current action or randomization, and is revealed only after u_t is chosen.
- 4) **Cost regularity:** each c_t is convex in (x, u) , obeys (5) on $\mathcal{D}_\alpha$, and admits the bounded subgradient oracle described above.
- 5) **Feasible policy class:** both the online policy and the benchmark policies are required to satisfy (1) under (3).

Let $\{(x_t, u_t)\}_{t=1}^T$ denote the realized trajectory of the online controller. Define its cumulative cost by

$$J_T^{\text{on}} \triangleq \sum_{t=1}^T c_t(x_t, u_t). \quad (6)$$

For a fixed positive integer H , let $\mathbf{m}$ range over the SDAC parameter set $\mathcal{M}_H$ (defined in Section IV). For each $\mathbf{m} \in \mathcal{M}_H$, write

$$J_T(\mathbf{m}) \triangleq \sum_{t=1}^T c_t(x_t(\mathbf{m}), u_t(\mathbf{m}))$$

for the cumulative cost of the induced feasible trajectory under the same arrival sequence. The regret against the best fixed SDAC policy is defined by

$$\text{Reg}_T^{\text{SDAC}} \triangleq J_T^{\text{on}} - \min_{\mathbf{m} \in \mathcal{M}_H} J_T(\mathbf{m}). \quad (7)$$

While the online controller may update its SDAC parameters over time, the benchmark in (7) is the best single fixed SDAC policy chosen in hindsight. The objective is to design an online policy that satisfies (1) for all $t \in [T]$ and achieves sublinear regret as defined in (7).

IV. SIMPLEX DISTURBANCE-ACTION CONTROL (SDAC)

A. From Disturbance-Action Control to SDAC

DAC combines a stabilizing state-feedback term with a finite linear filter of past disturbances. Under standard conditions, this parameterization yields convex surrogate losses and efficient online updates. In our scalar storage model, specializing the baseline feedback term to zero gives the representative form

$$u_t = \sum_{i=1}^H \theta_i w_{t-i}, \quad (8)$$

where H is a positive integer (the memory horizon) and $\boldsymbol{\theta} = (\theta_1, \dots, \theta_H)^\top \in \mathbb{R}^H$ is a linear filter.

For our constrained model, an unrestricted filter is insufficient: its action can violate (1) by exceeding αx_t or becoming negative. We therefore introduce *Simplex Disturbance-Action Control* (SDAC), which uses a simplex-constrained filter tailored to the storage dynamics. This restriction ensures that every fixed policy is feasible while preserving convex surrogate losses and tractable online updates.

¹We focus on $\alpha \in (0, 1)$; the case $\alpha \geq 1$ is discussed in Section VII.

B. SDAC Definition

Fix a memory horizon $H \in \mathbb{N}_{>0}$. Let

$$\mathcal{M}_H \triangleq \{\mathbf{m} \in \mathbb{R}_{\geq 0}^H : \|\mathbf{m}\|_1 \leq 1\} \quad (9)$$

be the subprobability simplex. For any $\mathbf{m} \in \mathcal{M}_H$, define the policy $\pi_{\mathbf{m}}$ by

$$u_t(\mathbf{m}) \triangleq \sum_{i=1}^H m_i \alpha^i w_{t-i} = \langle \boldsymbol{\phi}_t^u, \mathbf{m} \rangle, \quad (10)$$

where

$$\boldsymbol{\phi}_t^u \triangleq (\alpha w_{t-1}, \alpha^2 w_{t-2}, \dots, \alpha^H w_{t-H})^\top \in \mathbb{R}^H, \quad (11)$$

The induced state trajectory satisfies

$$x_{t+1}(\mathbf{m}) = \alpha x_t(\mathbf{m}) - u_t(\mathbf{m}) + w_t, \quad x_1(\mathbf{m}) = 0. \quad (12)$$

The SDAC policy class is

$$\Pi_H^{\text{SDAC}} \triangleq \{\pi_{\mathbf{m}} : \mathbf{m} \in \mathcal{M}_H\}. \quad (13)$$

The weights α^i align the arrival filter with the retention coefficient α in (3).

We now establish the two structural properties that make SDAC useful: the induced control action always satisfies the feasibility constraint (1), and the induced cost is convex and Lipschitz in $\mathbf{m}$, so that regret minimization against the best fixed SDAC policy is an instance of online convex optimization over the simplex. Both properties follow from an explicit decomposition of the state trajectory that isolates its dependence on $\mathbf{m}$.

Lemma 1 (State decomposition under SDAC). *Fix a positive integer H and $\mathbf{m} \in \mathcal{M}_H$. For $i \in \{1, \dots, H\}$, define*

$$r_{t,i} \triangleq \sum_{j=1}^i \alpha^{j-1} w_{t-j}, \quad r_t \triangleq \sum_{j=1}^{t-1} \alpha^{j-1} w_{t-j},$$

where the second sum is empty when $t = 1$. Then, for every $t \in [T]$,

$$x_t(\mathbf{m}) = \sum_{i=1}^H m_i r_{t,i} + (1 - \mathbf{1}^\top \mathbf{m}) r_t. \quad (14)$$

Define the quantity

$$\boldsymbol{\phi}_t^x \triangleq (r_{t,1} - r_t, \dots, r_{t,H} - r_t)^\top.$$

Then

$$x_t(\mathbf{m}) = r_t + \langle \boldsymbol{\phi}_t^x, \mathbf{m} \rangle.$$

Moreover,

$$0 \leq r_{t,i} \leq r_t \leq \frac{1}{1 - \alpha},$$

and $\mathbf{m} \mapsto x_t(\mathbf{m})$ is affine.

Proof. The auxiliary quantities satisfy

$$r_{t+1} = \alpha r_t + w_t, \quad r_1 = 0, \quad (15)$$

and, for every $i \in \{1, \dots, H\}$,

$$\begin{aligned} \alpha r_{t,i} - \alpha^i w_{t-i} &= \sum_{j=1}^{i-1} \alpha^j w_{t-j} \\ &= r_{t+1,i} - w_t. \end{aligned} \quad (16)$$

For $i = 1$, the sum in the first line is empty.

At $t = 1$, since $w_\tau \equiv 0$ for $\tau \leq 0$,

$$\begin{aligned} r_{1,i} &= \sum_{j=1}^i \alpha^{j-1} w_{1-j} = 0, \\ r_1 &= 0, \\ x_1(\mathbf{m}) &= 0, \end{aligned}$$

so (14) holds. Suppose it holds at time t . Using (12) and (10),

$$\begin{aligned} x_{t+1}(\mathbf{m}) &= \alpha \left(\sum_{i=1}^H m_i r_{t,i} + (1 - \mathbf{1}^\top \mathbf{m}) r_t \right) \\ &\quad - \sum_{i=1}^H m_i \alpha^i w_{t-i} + w_t \\ &= \sum_{i=1}^H m_i (\alpha r_{t,i} - \alpha^i w_{t-i}) \\ &\quad + \alpha (1 - \mathbf{1}^\top \mathbf{m}) r_t + w_t. \end{aligned} \quad (17)$$

Substituting (15) and (16) into (17),

$$\begin{aligned} x_{t+1}(\mathbf{m}) &= \sum_{i=1}^H m_i (r_{t+1,i} - w_t) \\ &\quad + (1 - \mathbf{1}^\top \mathbf{m}) (r_{t+1} - w_t) + w_t \\ &= \sum_{i=1}^H m_i r_{t+1,i} + (1 - \mathbf{1}^\top \mathbf{m}) r_{t+1}, \end{aligned} \quad (18)$$

because the total coefficient of w_t in the first line is $-\mathbf{1}^\top \mathbf{m} - (1 - \mathbf{1}^\top \mathbf{m}) + 1 = 0$. This proves (14) by induction.

All summands defining $r_{t,i}$ and r_t are nonnegative. Since $w_\tau = 0$ for $\tau \leq 0$, $r_{t,i}$ is a truncation of r_t ; hence

$$0 \leq r_{t,i} \leq r_t \leq \sum_{j=1}^{\infty} \alpha^{j-1} = \frac{1}{1-\alpha}.$$

Finally,

$$x_t(\mathbf{m}) = r_t + \sum_{i=1}^H m_i (r_{t,i} - r_t) = r_t + \langle \phi_t^x, \mathbf{m} \rangle,$$

which is affine in $\mathbf{m}$. $\square$

The first consequence of this decomposition is that every fixed policy is feasible.

Lemma 2 (Feasibility of SDAC policies). *For every $\mathbf{m} \in \mathcal{M}_H$, the trajectory induced by (10) and (12) satisfies the action constraint (1); explicitly,*

$$0 \leq u_t(\mathbf{m}) \leq \alpha x_t(\mathbf{m}) \quad \text{for all } t \in [T]. \quad (19)$$

Proof. Because m_i , α , and w_{t-i} are nonnegative,

$$u_t(\mathbf{m}) = \sum_{i=1}^H m_i \alpha^i w_{t-i} \geq 0.$$

For every $i \in \{1, \dots, H\}$,

$$\alpha r_{t,i} = \sum_{j=1}^i \alpha^j w_{t-j} \geq \alpha^i w_{t-i}. \quad (20)$$

Because $\mathbf{m} \in \mathcal{M}_H$ and $r_t \geq 0$ by Lemma 1,

$$\begin{aligned} \alpha x_t(\mathbf{m}) &= \sum_{i=1}^H m_i \alpha r_{t,i} + \alpha (1 - \mathbf{1}^\top \mathbf{m}) r_t \\ &\geq \sum_{i=1}^H m_i \alpha r_{t,i} \\ &\geq \sum_{i=1}^H m_i \alpha^i w_{t-i} = u_t(\mathbf{m}), \end{aligned} \quad (21)$$

where the second inequality uses (20). $\square$

For fixed arrivals, both coordinates of the SDAC state-action pair are affine in $\mathbf{m}$. The following lemma records the resulting convexity and the constants used by the online update.

Lemma 3 (Convexity and Lipschitz continuity of surrogate cost). *Define*

$$\ell_t(\mathbf{m}) \triangleq c_t(x_t(\mathbf{m}), u_t(\mathbf{m})), \quad \mathbf{m} \in \mathcal{M}_H,$$

and

$$G_\alpha \triangleq \alpha \left(\frac{L_x}{1-\alpha} + L_u \right). \quad (22)$$

Then ℓ_t is convex on $\mathcal{M}_H$, and for all $\mathbf{m}, \mathbf{m}' \in \mathcal{M}_H$,

$$|\ell_t(\mathbf{m}) - \ell_t(\mathbf{m}')| \leq G_\alpha \|\mathbf{m} - \mathbf{m}'\|_1.$$

Moreover, at every $\mathbf{m} \in \mathcal{M}_H$, one can select $\mathbf{g}_t \in \partial \ell_t(\mathbf{m})$ such that $\|\mathbf{g}_t\|_\infty \leq G_\alpha$.

Proof. Lemma 1 and (10) give

$$x_t(\mathbf{m}) = r_t + \langle \phi_t^x, \mathbf{m} \rangle, \quad u_t(\mathbf{m}) = \langle \phi_t^u, \mathbf{m} \rangle. \quad (23)$$

For every $i \in \{1, \dots, H\}$, the definitions of $r_{t,i}$ and r_t give

$$\begin{aligned} (\phi_t^x)_i &= r_{t,i} - r_t \\ &= - \sum_{j=i+1}^{t-1} \alpha^{j-1} w_{t-j}, \end{aligned} \quad (24)$$

where the sum is empty when $i \geq t-1$. Hence

$$\begin{aligned} |(\phi_t^x)_i| &\leq \sum_{j=i+1}^{\infty} \alpha^{j-1} = \frac{\alpha^i}{1-\alpha}, \\ 0 \leq (\phi_t^u)_i &= \alpha^i w_{t-i} \leq \alpha^i. \end{aligned} \quad (25)$$

Taking the largest coordinate, attained in these upper bounds at $i = 1$, yields

$$\|\phi_t^x\|_\infty \leq \frac{\alpha}{1-\alpha}, \quad \|\phi_t^u\|_\infty \leq \alpha. \quad (26)$$

Since (23) is affine in $\mathbf{m}$ and c_t is convex, ℓ_t is convex.

By (5) and (23),

$$\begin{aligned} |\ell_t(\mathbf{m}) - \ell_t(\mathbf{m}')| &\leq L_x |\langle \phi_t^x, \mathbf{m} - \mathbf{m}' \rangle| + L_u |\langle \phi_t^u, \mathbf{m} - \mathbf{m}' \rangle| \\ &\leq (L_x \|\phi_t^x\|_\infty + L_u \|\phi_t^u\|_\infty) \|\mathbf{m} - \mathbf{m}'\|_1 \\ &\leq \alpha \left(\frac{L_x}{1-\alpha} + L_u \right) \|\mathbf{m} - \mathbf{m}'\|_1 \\ &= G_\alpha \|\mathbf{m} - \mathbf{m}'\|_1. \end{aligned} \quad (27)$$

Let (ζ_t^x, ζ_t^u) denote the components of the bounded subgradient $\zeta_t(x_t(\mathbf{m}), u_t(\mathbf{m}))$ from Section III, and set

$$\mathbf{g}_t \triangleq \zeta_t^x \phi_t^x + \zeta_t^u \phi_t^u. \quad (28)$$

For every $\mathbf{m}' \in \mathcal{M}_H$, the subgradient inequality and (23) yield

$$\begin{aligned} \ell_t(\mathbf{m}') &\geq \ell_t(\mathbf{m}) + \zeta_t^x \langle \phi_t^x, \mathbf{m}' - \mathbf{m} \rangle + \zeta_t^u \langle \phi_t^u, \mathbf{m}' - \mathbf{m} \rangle \\ &= \ell_t(\mathbf{m}) + \langle \mathbf{g}_t, \mathbf{m}' - \mathbf{m} \rangle. \end{aligned}$$

Thus $\mathbf{g}_t \in \partial \ell_t(\mathbf{m})$. For each $i \in \{1, \dots, H\}$, the exact coordinate decay in (25) gives

$$\begin{aligned} |g_{t,i}| &\leq |\zeta_t^x| |(\phi_t^x)_i| + |\zeta_t^u| |(\phi_t^u)_i| \\ &\leq \alpha^i \left(\frac{L_x}{1-\alpha} + L_u \right) \\ &\leq \alpha \left(\frac{L_x}{1-\alpha} + L_u \right) = G_\alpha. \end{aligned} \quad (29)$$

Taking the maximum over i proves

$$\|\mathbf{g}_t\|_\infty \leq G_\alpha. \quad (30)$$

□

Regret minimization against the best fixed SDAC policy thus becomes online convex optimization over the simplex.

C. Expressiveness: Decreasing-Allocation Policies

Energy-harvesting policies for Bernoulli recharges can be described by the sequence of expenditures following each recharge [11], [12]. In the corresponding decreasing-allocation policy class, each unit arriving at time s is assigned the same decreasing allocation sequence, beginning at time $s+1$ to respect the information pattern. The sequence may have an arbitrary decreasing shape; geometric sequences induce proportional state-feedback policies.

Define the set of feasible decreasing allocation sequences

$$\mathcal{Q}_\alpha^\downarrow \triangleq \left\{ \mathbf{q} = (q_i)_{i \geq 1} : q_1 \geq q_2 \geq \dots \geq 0, \sum_{i=1}^{\infty} \frac{q_i}{\alpha^i} \leq 1 \right\}. \quad (31)$$

For $\mathbf{q} \in \mathcal{Q}_\alpha^\downarrow$, define

$$u_t(\mathbf{q}) \triangleq \sum_{i=1}^{t-1} q_i w_{t-i}, \quad (32)$$

$$x_{t+1}(\mathbf{q}) = \alpha x_t(\mathbf{q}) - u_t(\mathbf{q}) + w_t, \quad x_1(\mathbf{q}) = 0.$$

Here $q_i w_s$ is the amount from arrival w_s allocated i periods later. The weighted budget in (31) accounts for leakage: reserving q_i units for expenditure i periods after an arrival requires q_i/α^i units at the arrival time. Let

$$\Pi_\alpha^{\text{alloc}} \triangleq \{\pi_{\mathbf{q}} : \mathbf{q} \in \mathcal{Q}_\alpha^\downarrow\}.$$

Lemma 4 (Finite-memory approximation of decreasing-allocation policies). *Every policy in $\Pi_\alpha^{\text{alloc}}$ is feasible. Fix $\mathbf{q} \in \mathcal{Q}_\alpha^\downarrow$ and $H \geq 1$, and define*

$$m_{H,i}(\mathbf{q}) \triangleq \frac{q_i}{\alpha^i}, \quad i \in [H]. \quad (33)$$

Then $\mathbf{m}_H(\mathbf{q}) \in \mathcal{M}_H$. With

$$\text{Tail}_H(\mathbf{q}) \triangleq \sum_{i=H+1}^{\infty} q_i,$$

the corresponding trajectories satisfy, for every arrival sequence in $[0, 1]$,

$$0 \leq u_t(\mathbf{q}) - u_t(\mathbf{m}_H(\mathbf{q})) \leq \text{Tail}_H(\mathbf{q}), \quad (34)$$

$$0 \leq x_t(\mathbf{m}_H(\mathbf{q})) - x_t(\mathbf{q}) \leq \frac{\text{Tail}_H(\mathbf{q})}{1-\alpha}. \quad (35)$$

If each c_t satisfies (5), then

$$\begin{aligned} \left| \sum_{t=1}^T c_t(x_t(\mathbf{m}_H(\mathbf{q})), u_t(\mathbf{m}_H(\mathbf{q}))) - \sum_{t=1}^T c_t(x_t(\mathbf{q}), u_t(\mathbf{q})) \right| \\ \leq \left(L_u + \frac{L_x}{1-\alpha} \right) \text{Tail}_H(\mathbf{q}) T \\ \leq \left(L_u + \frac{L_x}{1-\alpha} \right) \alpha^{H+1} T. \end{aligned} \quad (36)$$

In particular, fix $k \in [0, \alpha]$, set $\lambda_k \triangleq \alpha - k$, and define

$$q_i(k) = k \lambda_k^{i-1}, \quad i \geq 1, \quad (37)$$

where $\lambda_k^0 \triangleq 1$. Then $\mathbf{q}(k) \in \mathcal{Q}_\alpha^\downarrow$ and $\pi_{\mathbf{q}(k)}$ is exactly the feasible proportional state-feedback policy $u_t(k) = k x_t(k)$. Its SDAC truncation is

$$[\mathbf{m}_H^{\text{prop}}(k)]_i \triangleq m_{H,i}(\mathbf{q}(k)) = \frac{k \lambda_k^{i-1}}{\alpha^i}, \quad i \in [H],$$

and its tail is

$$\text{Tail}_H(\mathbf{q}(k)) = \frac{k \lambda_k^H}{1-\lambda_k}.$$

Therefore, (36) specializes exactly to

$$\begin{aligned} \left| \sum_{t=1}^T c_t(x_t(\mathbf{m}_H^{\text{prop}}(k)), u_t(\mathbf{m}_H^{\text{prop}}(k))) - \sum_{t=1}^T c_t(x_t(k), u_t(k)) \right| \\ \leq \left(L_u + \frac{L_x}{1-\alpha} \right) \frac{k \lambda_k^H}{1-\lambda_k} T. \end{aligned} \quad (38)$$

The proof is deferred to Appendix A.

$\Pi_\alpha^{\text{alloc}}$ contains every feasible proportional state-feedback policy and also permits nongeometric allocation sequences. Together with Theorem 1, the tail bound in (36) yields no regret against $\Pi_\alpha^{\text{alloc}}$ and, consequently, against feasible proportional state-feedback policies.

V. ONLINE ALGORITHM

We now present our online algorithm, which updates SDAC parameters by OMD with unnormalized negentropy (equivalently, the exponentiated gradient method) and a fixed step size $\eta > 0$ over the subprobability simplex $\mathcal{M}_H \subset \mathbb{R}^H$ [26]. Let

$$\Psi(\mathbf{m}) \triangleq \sum_{i=1}^H (m_i \log m_i - m_i), \quad \mathbf{m} \in \mathbb{R}_{\geq 0}^H, \quad (39)$$

denote the unnormalized negentropy (with the convention $0 \log 0 \triangleq 0$). Its Bregman divergence is the generalized Kullback–Leibler divergence

$$D_\Psi(\mathbf{m} \parallel \mathbf{v}) \triangleq \sum_{i=1}^H \left(m_i \log \frac{m_i}{v_i} - m_i + v_i \right), \quad (40)$$

defined for

$$\mathbf{m} \in \mathbb{R}_{\geq 0}^H, \quad \mathbf{v} \in \mathbb{R}_{> 0}^H.$$

For each time $t \in [T]$, let $\mathbf{m}_t \in \mathcal{M}_H$ be the online parameter. We initialize at

$$\mathbf{m}_1 \triangleq \frac{1}{H+1} \mathbf{1} \quad (41)$$

so that $\sup_{\mathbf{m} \in \mathcal{M}_H} D_\Psi(\mathbf{m} \parallel \mathbf{m}_1) \leq \log(H+1)$ (see the proof of Theorem 1).

a) *Control-action selection:* The controller first forms the nominal SDAC action

$$\hat{u}_t \triangleq \langle \phi_t^u, \mathbf{m}_t \rangle, \quad (42)$$

and clips it to the admissible interval in (1):

$$u_t \triangleq \min\{\hat{u}_t, \alpha x_t\}. \quad (43)$$

The state then evolves according to (3). Because $\hat{u}_t \geq 0$, the clipping rule and induction on (3), starting from (2) with $w_t \in [0, 1]$, yield $x_t \geq 0$ for all $t \in [T+1]$ and show that the online trajectory satisfies (1).

The scalar correction in (43) also has an exact Bregman-projection equivalent.

Lemma 5 (Action clipping as a Bregman projection). *At time t , define the parameter set compatible with the available state by*

$$\mathcal{F}_t(x_t) \triangleq \{\mathbf{m} \in \mathcal{M}_H : \langle \phi_t^u, \mathbf{m} \rangle \leq \alpha x_t\}. \quad (44)$$

For the negentropy Ψ in (39), let

$$\mathbf{m}_t^{\text{exec}} \in \arg \min_{\mathbf{m} \in \mathcal{F}_t(x_t)} D_\Psi(\mathbf{m} \parallel \mathbf{m}_t). \quad (45)$$

Then

$$\langle \phi_t^u, \mathbf{m}_t^{\text{exec}} \rangle = \min\{\langle \phi_t^u, \mathbf{m}_t \rangle, \alpha x_t\} = u_t. \quad (46)$$

Thus the projection formulation and the scalar clipping rule induce exactly the same executed action.

Proof. The set $\mathcal{F}_t(x_t)$ is nonempty because it contains $\mathbf{0}$, and it is compact and convex. Since $\mathbf{m}_t$ is strictly positive coordinatewise, $D_\Psi(\cdot \parallel \mathbf{m}_t)$ is strictly convex on $\mathcal{M}_H$ and is minimized uniquely at $\mathbf{m}_t$ over any feasible set containing $\mathbf{m}_t$. Consequently, if

$$\langle \phi_t^u, \mathbf{m}_t \rangle \leq \alpha x_t,$$

then $\mathbf{m}_t^{\text{exec}} = \mathbf{m}_t$, which proves (46) in the feasible case.

Suppose instead that $\mathbf{m}_t$ is infeasible. We claim that the action constraint is active at $\mathbf{m}_t^{\text{exec}}$. If it were slack, then

$$\langle \phi_t^u, \mathbf{m}_t^{\text{exec}} \rangle < \alpha x_t < \langle \phi_t^u, \mathbf{m}_t \rangle. \quad (47)$$

For sufficiently small $\lambda \in (0, 1)$, the point

$$\mathbf{m}_\lambda \triangleq (1 - \lambda)\mathbf{m}_t^{\text{exec}} + \lambda\mathbf{m}_t$$

would still belong to $\mathcal{F}_t(x_t)$. Strict convexity and $D_\Psi(\mathbf{m}_t \parallel \mathbf{m}_t) = 0$ would then give

$$\begin{aligned} D_\Psi(\mathbf{m}_\lambda \parallel \mathbf{m}_t) &< (1 - \lambda)D_\Psi(\mathbf{m}_t^{\text{exec}} \parallel \mathbf{m}_t) \\ &< D_\Psi(\mathbf{m}_t^{\text{exec}} \parallel \mathbf{m}_t), \end{aligned}$$

contradicting (45). Hence

$$\langle \phi_t^u, \mathbf{m}_t^{\text{exec}} \rangle = \alpha x_t$$

in the infeasible case, which completes the proof. $\square$

The implemented rule remains the scalar minimum (43); the full H -dimensional vector $\mathbf{m}_t^{\text{exec}}$ is never calculated because only its scalar induced action is executed. Forming the feature vectors and their inner products, followed by the OMD update below, takes $O(H)$ time and $O(H)$ memory per round. In particular, the next unconstrained OMD update uses the internal iterate $\mathbf{m}_t$, not $\mathbf{m}_t^{\text{exec}}$. Therefore the internal iterates remain on the fixed domain $\mathcal{M}_H$, and the analysis does not acquire a time-varying comparator set. The direct implementation and its projection interpretation produce the same actions, states, and internal OMD iterates.

b) *Surrogate cost and OMD update:* After (w_t, c_t) is revealed, select a subgradient of the convex surrogate loss defined in Lemma 3, $\mathbf{g}_t \in \partial \ell_t(\mathbf{m}_t)$ satisfying $\|\mathbf{g}_t\|_\infty \leq G_\alpha$ as in (22), whose existence and construction are given in that lemma. This step uses the revealed cost function and the full-information subgradient oracle from Section III. We define the OMD update

$$\mathbf{m}_{t+1} \triangleq \arg \min_{\mathbf{m} \in \mathcal{M}_H} \left\{ \eta \langle \mathbf{g}_t, \mathbf{m} \rangle + D_\Psi(\mathbf{m} \parallel \mathbf{m}_t) \right\}. \quad (48)$$

c) *Closed-form solution:* Although (48) is a constrained optimization problem, it is easy to solve analytically because Ψ is separable and the objective is linear in $\mathbf{g}_t$. As shown in Lemma 7, the minimizer in (48) is exactly the result of two elementary steps. First a multiplicative step,

$$\tilde{\mathbf{m}}_{t+1,i} \triangleq m_{t,i} \exp(-\eta g_{t,i}), \quad i = 1, \dots, H, \quad (49)$$

where $g_{t,i}$ is the i -th coordinate of $\mathbf{g}_t$. This is followed by the Bregman projection onto $\mathcal{M}_H$,

$$\mathbf{m}_{t+1} \triangleq \text{Proj}_{\mathcal{M}_H}^\Psi(\tilde{\mathbf{m}}_{t+1}) = \frac{\tilde{\mathbf{m}}_{t+1}}{\max\{1, \|\tilde{\mathbf{m}}_{t+1}\|_1\}}, \quad (50)$$

where the projection operator is defined and derived in Lemma 6; the scaling is nontrivial only when $\|\tilde{\mathbf{m}}_{t+1}\|_1 > 1$.

VI. REGRET ANALYSIS

The following theorem states that the online controller achieves sublinear regret against every fixed SDAC policy $\pi_{\mathbf{m}}$, where $\mathbf{m}$ is any parameter in $\mathcal{M}_H$. The regret proof has two parts: a standard online convex optimization bound for the surrogate costs ℓ_t , and a stability bound showing that the online trajectory stays close to the SDAC policy trajectory driven by the same parameters. Some technical lemmas required for the proof are deferred to Appendix A.

Theorem 1 (Regret bound). *Suppose the cost functions satisfy the regularity and bounded-subgradient assumptions of Section III. Then the OMD controller with action clipping from*

Section V, initialized as in (41) and using step size $\eta > 0$, satisfies, for the regret defined in (7),

$$\text{Reg}_T^{\text{SDAC}} \leq \frac{\log(H+1)}{\eta} + \eta T \Gamma_\alpha, \quad (51)$$

where

$$\Gamma_\alpha \triangleq \frac{G_\alpha^2}{2} + \frac{\alpha G_\alpha (L_x + \alpha L_u)}{(1-\alpha)^2}, \quad (52)$$

and $G_\alpha = \alpha(\frac{L_x}{1-\alpha} + L_u)$ as in (22). In particular, the choice

$$\eta^* = \sqrt{\frac{\log(H+1)}{T \Gamma_\alpha}}$$

gives

$$\text{Reg}_T^{\text{SDAC}} \leq 2\sqrt{\Gamma_\alpha T \log(H+1)}.$$

For fixed $L_x, L_u > 0$, the leading factor obeys

$$\sqrt{\Gamma_\alpha} = \Theta(\alpha) \quad \text{as } \alpha \downarrow 0,$$

and

$$\sqrt{\Gamma_\alpha} = \Theta((1-\alpha)^{-3/2}) \quad \text{as } \alpha \uparrow 1.$$

The next result shows that the $\sqrt{T}$ dependence is unavoidable and quantifies the dependence on $1-\alpha$ that any causal feasible controller must incur. Its proof is deferred to Appendix A.

Theorem 2 (Minimax SDAC regret lower bound). *Fix $H \geq 1$ and $T \geq 2$. For every causal feasible online controller $\mathcal{A}$, possibly randomized, there exists a deterministic oblivious sequence of arrivals and costs satisfying the assumptions of Section III such that*

$$\mathbb{E}_{\mathcal{A}}[\text{Reg}_T^{\text{SDAC}}] \geq \frac{1}{2\sqrt{2}} \sqrt{\sum_{t=2}^T \beta_t^2}, \quad (53)$$

where

$$\beta_t \triangleq \alpha \left(L_u + \frac{L_x(1-\alpha^{t-2})}{1-\alpha} \right). \quad (54)$$

The expectation is only over the possible internal randomization of $\mathcal{A}$; the adversarial sequence whose existence is asserted is deterministic and fixed before the interaction.

Corollary 1 (Lower bound for sufficiently long horizons). *Define*

$$t_{\text{mix}}(\alpha) \triangleq \left\lceil \frac{\log 2}{\log(1/\alpha)} \right\rceil. \quad (55)$$

If $T \geq 2t_{\text{mix}}(\alpha) + 2$, then every causal feasible controller $\mathcal{A}$ admits a deterministic oblivious sequence for which

$$\mathbb{E}_{\mathcal{A}}[\text{Reg}_T^{\text{SDAC}}] \geq \frac{G_\alpha}{8} \sqrt{T}, \quad (56)$$

where $G_\alpha = \alpha(L_x/(1-\alpha) + L_u)$ as in (22).

The limit $\alpha \uparrow 1$ is more challenging because stored arrivals persist longer. In this limit, $t_{\text{mix}}(\alpha) = \Theta((1-\alpha)^{-1})$, and Corollary 1 gives

$$\mathbb{E}_{\mathcal{A}}[\text{Reg}_T^{\text{SDAC}}] = \Omega\left(\frac{\sqrt{T}}{1-\alpha}\right) \quad (57)$$

whenever $T \geq 2t_{\text{mix}}(\alpha) + 2$. The upper bound in Theorem 1 is $O((1-\alpha)^{-3/2} \sqrt{T \log(H+1)})$, so a factor $(1-\alpha)^{-1/2}$ between the two bounds remains open, apart from the memory-dependent logarithm.

The condition $T \geq 2t_{\text{mix}}(\alpha) + 2$ is essential for interpreting the large- α asymptotic. If T is fixed while $\alpha \uparrow 1$, then

$$\frac{1-\alpha^{t-2}}{1-\alpha} \rightarrow t-2,$$

and the right-hand side of (53) converges to the finite quantity

$$\frac{1}{2\sqrt{2}} \sqrt{\sum_{t=2}^T (L_u + L_x(t-2))^2}.$$

Thus the exact finite-horizon formula, rather than the long-horizon scaling in (57), applies in this regime.

The large- α dependence improves the coarser analysis that ignores the exact coordinate decay in (25) and the nonexpansiveness of $z \mapsto [z]_+$ used in (94). That analysis gives a leading factor of order $(1-\alpha)^{-2}$ as $\alpha \uparrow 1$. The sharper rate follows from the gradient and stability estimates.

Theorem 1 and Lemma 4 give the following comparator guarantee.

Corollary 2 (No-regret against decreasing-allocation policies). *Set*

$$H_T \triangleq \max \left\{ 1, \left\lceil \frac{\log T}{\log(1/\alpha)} \right\rceil \right\},$$

and run the OMD controller with action clipping, memory horizon $H = H_T$, and the step size η^* of Theorem 1. For $\mathbf{q} \in \mathcal{Q}_\alpha^l$, let

$$J_T(\mathbf{q}) \triangleq \sum_{t=1}^T c_t(x_t(\mathbf{q}), u_t(\mathbf{q}))$$

and

$$\text{Reg}_T^{\text{alloc}} \triangleq J_T^{\text{on}} - \inf_{\mathbf{q} \in \mathcal{Q}_\alpha^l} J_T(\mathbf{q}).$$

Then

$$\begin{aligned} \text{Reg}_T^{\text{alloc}} &\leq 2\sqrt{\Gamma_\alpha T \log(H_T + 1)} \\ &\quad + \left(L_u + \frac{L_x}{1-\alpha} \right) \alpha^{H_T+1} T. \end{aligned} \quad (58)$$

Here Γ_α is defined in (52). Since $\alpha^{H_T} \leq 1/T$ and $\log(H_T + 1) = O(\log \log T)$ as $T \rightarrow \infty$, the approximation term is $O(1)$. Consequently,

$$\text{Reg}_T^{\text{alloc}} = O\left(\sqrt{T \log \log T}\right) = o(T).$$

Proof. Fix $\mathbf{q} \in \mathcal{Q}_\alpha^l$ and use the SDAC comparator $\mathbf{m}_{H_T}(\mathbf{q})$ from Lemma 4. Theorem 1 and (36) give

$$\begin{aligned} J_T^{\text{on}} - J_T(\mathbf{q}) &\leq 2\sqrt{\Gamma_\alpha T \log(H_T + 1)} \\ &\quad + \left(L_u + \frac{L_x}{1-\alpha} \right) \text{Tail}_{H_T}(\mathbf{q}) T \\ &\leq 2\sqrt{\Gamma_\alpha T \log(H_T + 1)} \\ &\quad + \left(L_u + \frac{L_x}{1-\alpha} \right) \alpha^{H_T+1} T. \end{aligned}$$

The bound is uniform in $\mathbf{q}$, so taking the infimum comparator cost proves (58). Finally, $H_T \geq \log T / \log(1/\alpha)$ implies $\alpha^{H_T} \leq 1/T$, while $H_T = \Theta(\log T)$ implies $\log(H_T + 1) = O(\log \log T)$. $\square$

Because geometric allocation sequences induce proportional state feedback, the same controller also satisfies a gain-dependent bound for that policy family.

Corollary 3 (No-regret against proportional state-feedback policies). *Run the controller of Corollary 2. For $k \in [0, \alpha]$, let $J_T(k) \triangleq \sum_{t=1}^T c_t(x_t(k), u_t(k))$, where $u_t(k) = kx_t(k)$, and define*

$$\text{Reg}_T^{\text{prop}} \triangleq J_T^{\text{on}} - \min_{k \in [0, \alpha]} J_T(k).$$

For every $k \in [0, \alpha]$, with $\lambda_k = \alpha - k$,

$$J_T^{\text{on}} - J_T(k) \leq 2\sqrt{\Gamma_\alpha T \log(H_T + 1)} + \left(L_u + \frac{L_x}{1 - \alpha}\right) \frac{k\lambda_k^{H_T}}{1 - \lambda_k} T. \quad (59)$$

Moreover,

$$\begin{aligned} \text{Reg}_T^{\text{prop}} &\leq 2\sqrt{\Gamma_\alpha T \log(H_T + 1)} \\ &\quad + \left(L_u + \frac{L_x}{1 - \alpha}\right) \alpha^{H_T+1} T \\ &= O\left(\sqrt{T \log \log T}\right) = o(T). \end{aligned}$$

Proof. For fixed k , Theorem 1 applied to $\mathbf{m}_{H_T}^{\text{prop}}(k)$, followed by (38), gives (59). The geometric allocation sequence $\mathbf{q}(k)$ belongs to $\mathcal{Q}_\alpha^\downarrow$, so $\text{Tail}_{H_T}(\mathbf{q}(k)) \leq \alpha^{H_T+1}$ by Lemma 4. This makes the bound uniform in k ; minimizing over k proves the stated finite-time and asymptotic guarantees. $\square$

VII. EXTENSIONS

We briefly discuss three possible extensions. Establishing their full regret bounds is left for future work.

a) Retention $\alpha \geq 1$. For $\alpha \geq 1$, choose a fixed baseline drain Kx_t such that $\bar{\alpha} \triangleq \alpha - K \in (0, 1)$, and write $u_t = Kx_t + v_t$. The residual dynamics become $x_{t+1} = \bar{\alpha}x_t - v_t + w_t$. An SDAC filter for v_t should therefore use the powers $\bar{\alpha}^i$, not α^i . If $0 \leq v_t \leq \bar{\alpha}x_t$, then $0 \leq u_t \leq (K + \bar{\alpha})x_t = \alpha x_t$, so (1) holds. This reduction suggests an extension of the present analysis with transformed cost and Lipschitz constants; we do not pursue the details here.

b) Vector action / multi-task allocation. Consider a single storage state x_t with a vector action $\mathbf{u}_t \in \mathbb{R}_{\geq 0}^n$ that allocates resources across n tasks, subject to $\mathbf{1}^\top \mathbf{u}_t \leq \alpha x_t$, so that the total drain from storage is the sum of the task allocations. The dynamics remain $x_{t+1} = \alpha x_t - \mathbf{1}^\top \mathbf{u}_t + w_t$, and the cost $c_t(x_t, \mathbf{u}_t)$ is convex in $(x, \mathbf{u})$. The fixed-policy SDAC construction lifts by replacing the filter vector $\mathbf{m} \in \mathcal{M}_H$ with a matrix $\mathbf{M} \in \mathbb{R}_{\geq 0}^{n \times H}$ whose entries $M_{a,i}$ sum to at most one, with each row a corresponding to a task and each column i to a memory lag. OMD can then run over this set. If the nominal total allocation exceeds αx_t , the action vector is scaled by a nonnegative factor so that $\mathbf{1}^\top \mathbf{u}_t = \alpha x_t$.

c) Multiple storages with diagonal dynamics. Suppose the state is $\mathbf{x}_t \in \mathbb{R}_{\geq 0}^d$ and evolves as $\mathbf{x}_{t+1} = \mathbf{A}\mathbf{x}_t - \mathbf{u}_t + \mathbf{w}_t$ with $\mathbf{A} = \text{diag}(\alpha_1, \dots, \alpha_d)$, $\alpha_i \in (0, 1)$, arrivals $\mathbf{w}_t \in [0, 1]^d$, and per-coordinate constraints $0 \leq u_{t,i} \leq \alpha_i x_{t,i}$. The dynamics and feasible set then decouple by coordinate, and the natural SDAC class is a product of simplices. For a jointly convex cost, OMD operates jointly on this product set; the optimization separates by coordinate only when the cost is coordinate-separable. Cross-coupled storage systems (non-diagonal $\mathbf{A}$) require additional analysis.

These constructions retain the basic simplex/OMD structure, but their complete feasibility and regret analyses require separate treatment.

VIII. CONCLUSION

We studied adversarial online control of a storage system under a state-dependent action constraint. Every fixed SDAC policy is feasible by construction. For the time-varying online controller, the OMD update maintains the SDAC parameter constraint, while the clipping rule maintains feasibility of the actual trajectory. Exact feature decay and the contraction of the online state dynamics yield $O\left(\sqrt{T \log(H+1)}\right)$ regret against the best fixed SDAC policy, with upper-bound dependence $\Theta((1-\alpha)^{-3/2})$ as $\alpha \uparrow 1$. Once T is at least on the order of $(1-\alpha)^{-1}$, the minimax lower bound gives $\Omega((1-\alpha)^{-1}\sqrt{T})$ regret as $\alpha \uparrow 1$; hence a factor $(1-\alpha)^{-1/2}$ between the upper and lower bounds remains open. For fixed T and $\alpha \uparrow 1$, the exact finite-horizon lower bound must be used instead, because the long-horizon asymptotic does not apply. The SDAC approximation results further give $O(\sqrt{T \log \log T})$ regret against the best decreasing-allocation policy when $H = \Theta(\log T)$. This class contains all feasible proportional state-feedback policies and also permits nongeometric allocation sequences. The model applies directly to energy-harvesting battery management and related settings in which storage limits must be respected in real time. Section VII discusses possible extensions to $\alpha \geq 1$, vector allocation, and diagonal multi-storage systems.

APPENDIX A DEFERRED PROOFS

This appendix proves the minimax lower bound from Section VI and the expressiveness lemma from Section IV, together with the OMD and stability lemmas leading up to the upper regret bound.

We first prove the lower bound. Following the standard random-sign mechanism for online linear-optimization lower bounds [26], the construction embeds a two-comparator online linear problem into the two endpoint SDAC policies $\mathbf{0}$ and $\mathbf{e}_1$.

Proof of Theorem 2. Fix a causal feasible controller $\mathcal{A}$. Let $\xi_1, \dots, \xi_T$ be independent Rademacher random variables, sampled before the interaction and independently of the internal randomization of $\mathcal{A}$. Consider the oblivious random construction

$$w_t = 1, \quad c_t(x, u) = C_\alpha + \xi_t(L_x x - L_u u), \quad (60)$$

where

$$C_\alpha \triangleq \frac{L_x + \alpha L_u}{1 - \alpha}. \quad (61)$$

Each c_t is affine and hence convex. For any two points in $\mathcal{D}_\alpha$,

$$\begin{aligned} |c_t(x, u) - c_t(x', u')| &= |L_x(x - x') - L_u(u - u')| \\ &\leq L_x|x - x'| + L_u|u - u'|. \end{aligned} \quad (62)$$

The constant gradient $\xi_t(L_x, -L_u)$ also obeys the bounded-subgradient assumption. Moreover, for $(x, u) \in \mathcal{D}_\alpha$,

$$\begin{aligned} |L_x x - L_u u| &\leq L_x x + L_u u \\ &\leq \frac{L_x + \alpha L_u}{1 - \alpha} = C_\alpha. \end{aligned} \quad (63)$$

Thus the offset even makes the costs nonnegative on the feasible domain, although it will cancel from regret.

Because $H \geq 1$, both

$$\mathbf{m}^{(0)} \triangleq \mathbf{0}, \quad \mathbf{m}^{(1)} \triangleq \mathbf{e}_1 \quad (64)$$

belong to $\mathcal{M}_H$, where $\mathbf{e}_1$ is the first standard basis vector in $\mathbb{R}^H$. Under the arrivals in (60), the first endpoint has

$$u_t(\mathbf{m}^{(0)}) = 0, \quad x_t(\mathbf{m}^{(0)}) = \rho_t \triangleq \frac{1 - \alpha^{t-1}}{1 - \alpha}. \quad (65)$$

Indeed, this follows by unrolling $x_{t+1} = \alpha x_t + 1$ from $x_1 = 0$. At the second endpoint,

$$u_1(\mathbf{m}^{(1)}) = 0, \quad u_t(\mathbf{m}^{(1)}) = \alpha, \quad t \geq 2. \quad (66)$$

The state satisfies $x_2(\mathbf{m}^{(1)}) = 1$. If it equals one at a time $t \geq 2$, then

$$x_{t+1}(\mathbf{m}^{(1)}) = \alpha - \alpha + 1 = 1.$$

Consequently,

$$x_t(\mathbf{m}^{(1)}) = 1, \quad t \geq 2. \quad (67)$$

For $t \geq 2$, the state difference between the endpoints is

$$\begin{aligned} \Delta_t^{\text{end}} &\triangleq x_t(\mathbf{m}^{(0)}) - x_t(\mathbf{m}^{(1)}) \\ &= \rho_t - 1 \\ &= \frac{\alpha(1 - \alpha^{t-2})}{1 - \alpha}. \end{aligned} \quad (68)$$

Remove the common offset by defining the centered online and endpoint costs

$$\begin{aligned} J_{T,\text{ctr}}^{\text{on}} &\triangleq \sum_{t=1}^T \xi_t(L_x x_t - L_u u_t), \\ J_{T,\text{ctr}}^{(j)} &\triangleq \sum_{t=1}^T \xi_t \left(L_x x_t(\mathbf{m}^{(j)}) - L_u u_t(\mathbf{m}^{(j)}) \right), \quad j \in \{0, 1\}. \end{aligned} \quad (69)$$

At time t , the pair (x_t, u_t) is determined by the past observations and the controller randomization used before c_t is revealed. It is therefore independent of the fresh sign ξ_t . Conditional mean zero gives

$$\begin{aligned} \mathbb{E}_{\xi, \mathcal{A}}[\xi_t(L_x x_t - L_u u_t) | x_t, u_t] &= 0, \\ \mathbb{E}_{\xi, \mathcal{A}}[J_{T,\text{ctr}}^{\text{on}}] &= 0. \end{aligned} \quad (70)$$

Likewise, the two fixed endpoint trajectories do not depend on the signs, so

$$\mathbb{E}_{\xi}[J_{T,\text{ctr}}^{(0)}] = \mathbb{E}_{\xi}[J_{T,\text{ctr}}^{(1)}] = 0. \quad (71)$$

Since the two endpoints agree at $t = 1$, equations (65), (66), and (68) give

$$\begin{aligned} J_{T,\text{ctr}}^{(0)} - J_{T,\text{ctr}}^{(1)} &= \sum_{t=2}^T \xi_t(L_x \Delta_t^{\text{end}} + \alpha L_u) \\ &= \sum_{t=2}^T \xi_t \beta_t, \end{aligned} \quad (72)$$

with β_t as in (54).

Minimizing over all of $\mathcal{M}_H$ can only decrease the comparator cost. After canceling the common term $T C_\alpha$, the regret therefore satisfies pointwise

$$\text{Reg}_T^{\text{SDAC}} \geq J_{T,\text{ctr}}^{\text{on}} - \min\{J_{T,\text{ctr}}^{(0)}, J_{T,\text{ctr}}^{(1)}\}. \quad (73)$$

Using $\min\{a, b\} = (a + b - |a - b|)/2$ together with (70) and (71), we obtain

$$\begin{aligned} \mathbb{E}_{\xi, \mathcal{A}}[\text{Reg}_T^{\text{SDAC}}] &\geq \frac{1}{2} \mathbb{E}_{\xi} \left[J_{T,\text{ctr}}^{(0)} - J_{T,\text{ctr}}^{(1)} \right] \\ &= \frac{1}{2} \mathbb{E}_{\xi} \left[\sum_{t=2}^T \xi_t \beta_t \right]. \end{aligned} \quad (74)$$

Khintchine's inequality for Rademacher sums implies

$$\mathbb{E}_{\xi} \left[\left| \sum_{t=2}^T \xi_t \beta_t \right| \right] \geq \frac{1}{\sqrt{2}} \left(\sum_{t=2}^T \beta_t^2 \right)^{1/2}. \quad (75)$$

Substitution of (54) into (74)–(75) proves (53) under the joint randomization of the signs and the controller.

It remains to make the adversarial sequence deterministic. For a fixed sign vector $\xi \in \{-1, 1\}^T$, let

$$R_{\mathcal{A}}(\xi) \triangleq \mathbb{E}_{\mathcal{A}}[\text{Reg}_T^{\text{SDAC}} | \xi].$$

The bound just proved states that the average of $R_{\mathcal{A}}(\xi)$ over the 2^T sign vectors is at least the right-hand side of (53). Hence at least one deterministic vector ξ^* attains that value. Fixing ξ^* in (60) produces a deterministic oblivious admissible sequence, and the only remaining expectation is over the internal randomization of $\mathcal{A}$. $\square$

Proof of Corollary 1. The definition of $t_{\text{mix}}(\alpha)$ gives

$$\alpha^{t_{\text{mix}}(\alpha)} \leq \alpha^{\log 2 / \log(1/\alpha)} = \frac{1}{2}. \quad (76)$$

Therefore, for every $t \geq t_{\text{mix}}(\alpha) + 2$,

$$\begin{aligned} \beta_t &\geq \alpha \left(L_u + \frac{L_x}{2(1 - \alpha)} \right) \\ &\geq \frac{\alpha}{2} \left(L_u + \frac{L_x}{1 - \alpha} \right) = \frac{G_\alpha}{2}. \end{aligned} \quad (77)$$

There are $T - t_{\text{mix}}(\alpha) - 1$ indices in $\{t_{\text{mix}}(\alpha) + 2, \dots, T\}$. If $T \geq 2t_{\text{mix}}(\alpha) + 2$, then

$$T - t_{\text{mix}}(\alpha) - 1 \geq \frac{T}{2}. \quad (78)$$

Applying Theorem 2 and retaining only these indices,

$$\begin{aligned}\mathbb{E}_{\mathcal{A}}[\text{Reg}_{ST}^{\text{SDAC}}] &\geq \frac{1}{2\sqrt{2}} \left[\frac{T}{2} \left(\frac{G_{\alpha}}{2} \right)^2 \right]^{1/2} \\ &= \frac{G_{\alpha}}{8} \sqrt{T}.\end{aligned}\quad (79)$$

This is (56). $\square$

We next prove the expressiveness result stated in Section IV-C.

Proof of Lemma 4. Fix $\mathbf{q} \in \mathcal{Q}_{\alpha}^{\downarrow}$ and set

$$a_i \triangleq \frac{q_i}{\alpha^i}, \quad b_j \triangleq 1 - \sum_{i=1}^{j-1} a_i, \quad j \geq 1.$$

The weighted-budget constraint gives $a_i \geq 0$ and $b_j \geq 0$. Moreover,

$$x_t(\mathbf{q}) = \sum_{j=1}^{t-1} \alpha^{j-1} b_j w_{t-j}. \quad (80)$$

Indeed, the identity is immediate at $t = 1$. If it holds at time t , then

$$\begin{aligned}x_{t+1}(\mathbf{q}) &= w_t + \sum_{j=1}^{t-1} (\alpha^j b_j - q_j) w_{t-j} \\ &= w_t + \sum_{j=1}^{t-1} \alpha^j b_{j+1} w_{t-j},\end{aligned}$$

because $q_j = \alpha^j a_j$ and $b_{j+1} = b_j - a_j$. Reindexing gives (80) at time $t + 1$.

For each $j \geq 1$,

$$a_j \leq 1 - \sum_{i=1}^{j-1} a_i = b_j.$$

Consequently,

$$\begin{aligned}0 \leq u_t(\mathbf{q}) &= \sum_{j=1}^{t-1} \alpha^j a_j w_{t-j} \\ &\leq \sum_{j=1}^{t-1} \alpha^j b_j w_{t-j} = \alpha x_t(\mathbf{q}),\end{aligned}$$

which proves (1) for the decreasing-allocation policy.

For the truncated vector in (33),

$$m_{H,i}(\mathbf{q}) = a_i \geq 0, \quad \sum_{i=1}^H m_{H,i}(\mathbf{q}) \leq \sum_{i=1}^{\infty} a_i \leq 1.$$

Thus $\mathbf{m}_H(\mathbf{q}) \in \mathcal{M}_H$. Its SDAC action is

$$u_t(\mathbf{m}_H(\mathbf{q})) = \sum_{i=1}^{\min\{H, t-1\}} q_i w_{t-i}.$$

It follows that

$$0 \leq u_t(\mathbf{q}) - u_t(\mathbf{m}_H(\mathbf{q})) = \sum_{i=H+1}^{t-1} q_i w_{t-i} \leq \text{Tail}_H(\mathbf{q}),$$

where the middle sum is empty when $t \leq H + 1$. This proves (34).

Define

$$e_t^x \triangleq x_t(\mathbf{m}_H(\mathbf{q})) - x_t(\mathbf{q}), \quad e_t^u \triangleq u_t(\mathbf{q}) - u_t(\mathbf{m}_H(\mathbf{q})).$$

The two state recursions give

$$e_{t+1}^x = \alpha e_t^x + e_t^u, \quad e_1^x = 0.$$

Since $0 \leq e_t^u \leq \text{Tail}_H(\mathbf{q})$,

$$0 \leq e_t^x = \sum_{s=1}^{t-1} \alpha^{t-1-s} e_s^u \leq \frac{\text{Tail}_H(\mathbf{q})}{1-\alpha},$$

which proves (35). The Lipschitz property now yields

$$\begin{aligned}&|c_t(x_t(\mathbf{m}_H(\mathbf{q})), u_t(\mathbf{m}_H(\mathbf{q}))) \\ &\quad - c_t(x_t(\mathbf{q}), u_t(\mathbf{q}))| \\ &\leq \left(L_u + \frac{L_x}{1-\alpha} \right) \text{Tail}_H(\mathbf{q}).\end{aligned}$$

Finally,

$$\begin{aligned}\text{Tail}_H(\mathbf{q}) &= \sum_{i=H+1}^{\infty} \alpha^i a_i \\ &\leq \alpha^{H+1} \sum_{i=H+1}^{\infty} a_i \leq \alpha^{H+1}.\end{aligned}$$

Summing the per-period cost bound proves (36).

It remains to identify proportional state-feedback policies within the decreasing-allocation policy class. For $k \in [0, \alpha]$, let $\lambda_k = \alpha - k$ and let $\mathbf{q}(k)$ be defined by (37). The sequence is nonnegative and decreasing. If $k = 0$, then $\mathbf{q}(k) = \mathbf{0} \in \mathcal{Q}_{\alpha}^{\downarrow}$. If $k > 0$, then $\lambda_k/\alpha < 1$ and

$$\sum_{i=1}^{\infty} \frac{q_i(k)}{\alpha^i} = \frac{k}{\alpha} \sum_{j=0}^{\infty} \left(\frac{\lambda_k}{\alpha} \right)^j = \frac{k}{\alpha - \lambda_k} = 1,$$

so again $\mathbf{q}(k) \in \mathcal{Q}_{\alpha}^{\downarrow}$. Substituting $u_t(k) = kx_t(k)$ into the dynamics and unrolling gives

$$u_t(k) = k \sum_{i=1}^{t-1} \lambda_k^{i-1} w_{t-i} = \sum_{i=1}^{t-1} q_i(k) w_{t-i} = u_t(\mathbf{q}(k)).$$

The two policies therefore have identical action and state trajectories. Their SDAC truncations also coincide coordinate-wise, and the omitted geometric tail is

$$\text{Tail}_H(\mathbf{q}(k)) = k \sum_{i=H+1}^{\infty} \lambda_k^{i-1} = \frac{k \lambda_k^H}{1 - \lambda_k}.$$

Substitution in (36) proves (38). $\square$

The remaining lemmas establish the OMD and stability estimates used in the proof of Theorem 1; they rely on the state-decomposition, feasibility, and surrogate-cost lemmas already established in Section IV (Lemmas 1, 2, and 3).

Lemma 6 (Bregman projection onto $\mathcal{M}_H$). *Let D_{Ψ} be the Bregman divergence (40) of the unnormalized negentropy (39). Fix $\tilde{\mathbf{m}} \in \mathbb{R}_{>0}^H$. Then the Bregman projection of $\tilde{\mathbf{m}}$ onto $\mathcal{M}_H$,*

$$\text{Proj}_{\mathcal{M}_H}^{\Psi}(\tilde{\mathbf{m}}) \triangleq \arg \min_{\mathbf{m} \in \mathcal{M}_H} D_{\Psi}(\mathbf{m} \| \tilde{\mathbf{m}}),$$

is given by

$$\text{Proj}_{\mathcal{M}_H}^{\Psi}(\tilde{\mathbf{m}}) = \frac{\tilde{\mathbf{m}}}{\max\{1, \|\tilde{\mathbf{m}}\|_1\}}.$$

Proof. The projection objective $\mathbf{m} \mapsto D_{\Psi}(\mathbf{m} \|\tilde{\mathbf{m}})$ is strictly convex, and, for $m_i > 0$,

$$\frac{\partial}{\partial m_i} D_{\Psi}(\mathbf{m} \|\tilde{\mathbf{m}}) = \log \frac{m_i}{\tilde{m}_i}. \quad (81)$$

Its unconstrained minimizer is therefore $\mathbf{m} = \tilde{\mathbf{m}}$.

If $\|\tilde{\mathbf{m}}\|_1 \leq 1$, this minimizer belongs to $\mathcal{M}_H$, and hence

$$\text{Proj}_{\mathcal{M}_H}^{\Psi}(\tilde{\mathbf{m}}) = \tilde{\mathbf{m}}.$$

Suppose $\|\tilde{\mathbf{m}}\|_1 > 1$. The constraint $\|\mathbf{m}\|_1 \leq 1$ is then active. Moreover, the minimizer has no zero coordinate: if $m_{i_0} = 0$, transferring an arbitrarily small amount of mass from a positive coordinate to coordinate i_0 has directional derivative $-\infty$, because

$$\lim_{z \downarrow 0} \log \frac{z}{\tilde{m}_{i_0}} = -\infty.$$

Thus the nonnegativity constraints are inactive. With $\nu \geq 0$ denoting the multiplier of $\|\mathbf{m}\|_1 \leq 1$, stationarity and complementary slackness give

$$\log \frac{m_i}{\tilde{m}_i} + \nu = 0, \quad \|\mathbf{m}\|_1 = 1. \quad (82)$$

The first relation in (82) implies

$$m_i = \tilde{m}_i \exp(-\nu).$$

Summing over i and using the second relation in (82),

$$1 = \|\tilde{\mathbf{m}}\|_1 \exp(-\nu), \quad \exp(-\nu) = \frac{1}{\|\tilde{\mathbf{m}}\|_1}.$$

Hence $\nu = \log \|\tilde{\mathbf{m}}\|_1 > 0$ and

$$\text{Proj}_{\mathcal{M}_H}^{\Psi}(\tilde{\mathbf{m}}) = \frac{\tilde{\mathbf{m}}}{\|\tilde{\mathbf{m}}\|_1}.$$

Combining the cases $\|\tilde{\mathbf{m}}\|_1 \leq 1$ and $\|\tilde{\mathbf{m}}\|_1 > 1$ proves the stated formula. $\square$

Lemma 7 (Closed-form solution of the OMD update). *Fix $\mathbf{m}_t \in \mathcal{M}_H \cap \mathbb{R}_{>0}^H$, $\eta > 0$, and $\mathbf{g}_t \in \mathbb{R}^H$, and let $\tilde{\mathbf{m}}_{t+1}$ be defined coordinatewise by (49), i.e. $\tilde{m}_{t+1,i} = m_{t,i} \exp(-\eta g_{t,i})$. Then*

$$\begin{aligned} \arg \min_{\mathbf{m} \in \mathcal{M}_H} \{ \eta \langle \mathbf{g}_t, \mathbf{m} \rangle + D_{\Psi}(\mathbf{m} \|\mathbf{m}_t) \} \\ = \arg \min_{\mathbf{m} \in \mathcal{M}_H} D_{\Psi}(\mathbf{m} \|\tilde{\mathbf{m}}_{t+1}) = \text{Proj}_{\mathcal{M}_H}^{\Psi}(\tilde{\mathbf{m}}_{t+1}), \end{aligned}$$

where $\text{Proj}_{\mathcal{M}_H}^{\Psi}$ is the Bregman projection of Lemma 6. Consequently, the OMD update (48) coincides exactly with the closed-form OMD update (49)–(50), and in particular $\mathbf{m}_{t+1} \in \mathcal{M}_H \cap \mathbb{R}_{>0}^H$. Thus the positive initialization (41) keeps every algorithm iterate strictly positive.

Proof. Fix $\mathbf{m} \in \mathbb{R}_{>0}^H$ and a coordinate i . Then

$$\begin{aligned} \eta g_{t,i} m_i + m_i \log \frac{m_i}{m_{t,i}} &= m_i \left(\eta g_{t,i} + \log \frac{m_i}{m_{t,i}} \right) \\ &= m_i \log \frac{m_i}{\tilde{m}_{t+1,i}}. \end{aligned} \quad (83)$$

Substitution into the OMD objective gives

$$\begin{aligned} \eta \langle \mathbf{g}_t, \mathbf{m} \rangle + D_{\Psi}(\mathbf{m} \|\mathbf{m}_t) \\ = D_{\Psi}(\mathbf{m} \|\tilde{\mathbf{m}}_{t+1}) + \sum_{i=1}^H (m_{t,i} - \tilde{m}_{t+1,i}). \end{aligned} \quad (84)$$

The final sum is independent of $\mathbf{m}$. Hence the two optimization problems have the same minimizer, and Lemma 6 gives (50). Since $\mathbf{m}_t > 0$ coordinatewise, both the multiplicative step and the subsequent positive rescaling preserve strict positivity. $\square$

Lemma 8 (One-step movement of OMD). *For the OMD update with step size $\eta > 0$ and subgradient $\mathbf{g}_t \in \partial \ell_t(\mathbf{m}_t)$,*

$$\|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1 \leq \eta \|\mathbf{g}_t\|_{\infty}.$$

In particular, since $\|\mathbf{g}_t\|_{\infty} \leq G_{\alpha}$, $\|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1 \leq \eta G_{\alpha}$.

Proof. First-order optimality of (48), evaluated at the feasible point $\mathbf{m}_t$, gives

$$\begin{aligned} \langle \eta \mathbf{g}_t + \nabla \Psi(\mathbf{m}_{t+1}) - \nabla \Psi(\mathbf{m}_t), \mathbf{m}_t - \mathbf{m}_{t+1} \rangle &\geq 0, \\ D_{\Psi}(\mathbf{m}_{t+1} \|\mathbf{m}_t) + D_{\Psi}(\mathbf{m}_t \|\mathbf{m}_{t+1}) \\ &\leq \eta \langle \mathbf{g}_t, \mathbf{m}_t - \mathbf{m}_{t+1} \rangle \\ &\leq \eta \|\mathbf{g}_t\|_{\infty} \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1. \end{aligned} \quad (85)$$

The second line uses the symmetric Bregman identity.

For every $\mathbf{z} \in \mathcal{M}_H \cap \mathbb{R}_{>0}^H$ and $\mathbf{v} \in \mathbb{R}^H$,

$$\begin{aligned} \mathbf{v}^{\top} \nabla^2 \Psi(\mathbf{z}) \mathbf{v} &= \sum_{i=1}^H \frac{v_i^2}{z_i} \\ &\geq \frac{\|\mathbf{v}\|_1^2}{\sum_{i=1}^H z_i} \\ &\geq \|\mathbf{v}\|_1^2. \end{aligned} \quad (86)$$

The first inequality is Cauchy–Schwarz. The segment joining $\mathbf{m}_t$ and $\mathbf{m}_{t+1}$ lies in $\mathcal{M}_H \cap \mathbb{R}_{>0}^H$ by Lemma 7; integrating (86) along this segment yields

$$\begin{aligned} D_{\Psi}(\mathbf{m}_{t+1} \|\mathbf{m}_t) &\geq \frac{1}{2} \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1^2, \\ D_{\Psi}(\mathbf{m}_t \|\mathbf{m}_{t+1}) &\geq \frac{1}{2} \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1^2. \end{aligned}$$

Combining these inequalities with (85),

$$\|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1^2 \leq \eta \|\mathbf{g}_t\|_{\infty} \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1.$$

The claim is immediate if $\mathbf{m}_{t+1} = \mathbf{m}_t$; otherwise divide by $\|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1$. $\square$

Lemma 9 (OMD regret for the surrogate losses). *For every $\mathbf{m} \in \mathcal{M}_H$, the update (48) satisfies*

$$\sum_{t=1}^T \ell_t(\mathbf{m}_t) - \sum_{t=1}^T \ell_t(\mathbf{m}) \leq \frac{D_{\Psi}(\mathbf{m} \|\mathbf{m}_1)}{\eta} + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{g}_t\|_{\infty}^2. \quad (87)$$

Proof. Convexity of ℓ_t gives

$$\begin{aligned} \ell_t(\mathbf{m}_t) - \ell_t(\mathbf{m}) &\leq \langle \mathbf{g}_t, \mathbf{m}_t - \mathbf{m} \rangle \\ &= \langle \mathbf{g}_t, \mathbf{m}_t - \mathbf{m}_{t+1} \rangle \\ &\quad + \langle \mathbf{g}_t, \mathbf{m}_{t+1} - \mathbf{m} \rangle. \end{aligned} \quad (88)$$

First-order optimality of (48), evaluated at $\mathbf{m} \in \mathcal{M}_H$, yields

$$\begin{aligned} \eta \langle \mathbf{g}_t, \mathbf{m}_{t+1} - \mathbf{m} \rangle &\leq \langle \nabla \Psi(\mathbf{m}_{t+1}) - \nabla \Psi(\mathbf{m}_t), \mathbf{m} - \mathbf{m}_{t+1} \rangle \\ &= D_\Psi(\mathbf{m} \| \mathbf{m}_t) - D_\Psi(\mathbf{m} \| \mathbf{m}_{t+1}) \\ &\quad - D_\Psi(\mathbf{m}_{t+1} \| \mathbf{m}_t). \end{aligned} \quad (89)$$

The equality is the three-point Bregman identity. Substituting (89) into (88),

$$\begin{aligned} \ell_t(\mathbf{m}_t) - \ell_t(\mathbf{m}) &\leq \frac{D_\Psi(\mathbf{m} \| \mathbf{m}_t)}{\eta} - \frac{D_\Psi(\mathbf{m} \| \mathbf{m}_{t+1})}{\eta} \\ &\quad + \langle \mathbf{g}_t, \mathbf{m}_t - \mathbf{m}_{t+1} \rangle \\ &\quad - \frac{D_\Psi(\mathbf{m}_{t+1} \| \mathbf{m}_t)}{\eta}. \end{aligned} \quad (90)$$

The strong-convexity calculation (86) and Hölder's inequality give

$$\begin{aligned} \langle \mathbf{g}_t, \mathbf{m}_t - \mathbf{m}_{t+1} \rangle &- \frac{D_\Psi(\mathbf{m}_{t+1} \| \mathbf{m}_t)}{\eta} \\ &\leq \|\mathbf{g}_t\|_\infty \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1 - \frac{\|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1^2}{2\eta} \\ &\leq \frac{\eta}{2} \|\mathbf{g}_t\|_\infty^2. \end{aligned} \quad (91)$$

The last inequality is Young's inequality. Combining (90) and (91), and then summing over t , gives

$$\begin{aligned} \sum_{t=1}^T \ell_t(\mathbf{m}_t) - \sum_{t=1}^T \ell_t(\mathbf{m}) &\leq \frac{D_\Psi(\mathbf{m} \| \mathbf{m}_1)}{\eta} \\ &\quad - \frac{D_\Psi(\mathbf{m} \| \mathbf{m}_{T+1})}{\eta} \\ &\quad + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{g}_t\|_\infty^2. \end{aligned}$$

Since $D_\Psi(\mathbf{m} \| \mathbf{m}_{T+1}) \geq 0$, this proves (87). $\square$

Lemma 10 (Positive-part representation of the dynamics). *For $\mathbf{m} \in \mathcal{M}_H$, define*

$$F_t(x; \mathbf{m}) \triangleq [\alpha x - \langle \phi_t^u, \mathbf{m} \rangle]_+ + w_t. \quad (92)$$

Then the online state and every fixed-policy state satisfy

$$x_{t+1} = F_t(x_t; \mathbf{m}_t), \quad x_{t+1}(\mathbf{m}) = F_t(x_t(\mathbf{m}); \mathbf{m}). \quad (93)$$

Moreover, for all $x, x' \geq 0$ and $\mathbf{m}, \mathbf{m}' \in \mathcal{M}_H$,

$$\begin{aligned} |F_t(x; \mathbf{m}) - F_t(x'; \mathbf{m}')| &\leq |\alpha(x - x') - \langle \phi_t^u, \mathbf{m} - \mathbf{m}' \rangle| \\ &\leq \alpha|x - x'| + \|\phi_t^u\|_\infty \|\mathbf{m} - \mathbf{m}'\|_1. \end{aligned} \quad (94)$$

Proof. For the online trajectory, (43) implies

$$\begin{aligned} \alpha x_t - u_t &= \alpha x_t - \min\{\langle \phi_t^u, \mathbf{m}_t \rangle, \alpha x_t\} \\ &= [\alpha x_t - \langle \phi_t^u, \mathbf{m}_t \rangle]_+. \end{aligned} \quad (95)$$

For a fixed $\mathbf{m}$, Lemma 2 gives

$$\alpha x_t(\mathbf{m}) - \langle \phi_t^u, \mathbf{m} \rangle \geq 0,$$

so inserting a positive part does not change its state update. This proves (93). Finally, the map $z \mapsto [z]_+$ is nonexpansive:

$$|[z]_+ - [z']_+| \leq |z - z'|.$$

Apply this inequality to the two arguments in (92), and then apply Hölder's inequality, to obtain (94). $\square$

Lemma 11 (Contractive state and action deviation). *For the online trajectory and the counterfactual fixed-policy trajectory associated with the current internal iterate, one has, for every $t \in [T]$,*

$$|x_t - x_t(\mathbf{m}_t)| \leq \frac{\alpha \eta G_\alpha}{(1 - \alpha)^2}, \quad (96)$$

$$0 \leq u_t(\mathbf{m}_t) - u_t \leq \frac{\alpha^2 \eta G_\alpha}{(1 - \alpha)^2}. \quad (97)$$

Here $x_t(\mathbf{m}_t)$ is generated by using the single fixed parameter $\mathbf{m}_t$ from time 1 through time t .

Proof. Define

$$d_t \triangleq |x_t - x_t(\mathbf{m}_t)|.$$

Since $x_1 = x_1(\mathbf{m}_1) = 0$, one has $d_1 = 0$. For $t \in \{1, \dots, T-1\}$, Lemma 10 gives

$$\begin{aligned} d_{t+1} &= |F_t(x_t; \mathbf{m}_t) - F_t(x_t(\mathbf{m}_{t+1}); \mathbf{m}_{t+1})| \\ &\leq \alpha |x_t - x_t(\mathbf{m}_{t+1})| + \|\phi_t^u\|_\infty \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1. \end{aligned} \quad (98)$$

The affine state representation (23) implies

$$\begin{aligned} |x_t(\mathbf{m}_t) - x_t(\mathbf{m}_{t+1})| &= |\langle \phi_t^x, \mathbf{m}_t - \mathbf{m}_{t+1} \rangle| \\ &\leq \|\phi_t^x\|_\infty \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1. \end{aligned} \quad (99)$$

Combining the triangle inequality with (98) and (99) yields

$$\begin{aligned} d_{t+1} &\leq \alpha d_t \\ &\quad + (\alpha \|\phi_t^x\|_\infty + \|\phi_t^u\|_\infty) \|\mathbf{m}_{t+1} - \mathbf{m}_t\|_1. \end{aligned} \quad (100)$$

By (26),

$$\begin{aligned} \alpha \|\phi_t^x\|_\infty + \|\phi_t^u\|_\infty &\leq \frac{\alpha^2}{1 - \alpha} + \alpha \\ &= \frac{\alpha}{1 - \alpha}. \end{aligned} \quad (101)$$

Lemma 8 and $\|\mathbf{g}_t\|_\infty \leq G_\alpha$ therefore reduce (100) to

$$d_{t+1} \leq \alpha d_t + \frac{\alpha \eta G_\alpha}{1 - \alpha}. \quad (102)$$

Unrolling from $d_1 = 0$ gives

$$\begin{aligned} d_t &\leq \frac{\alpha \eta G_\alpha}{1 - \alpha} \sum_{j=0}^{t-2} \alpha^j \\ &\leq \frac{\alpha \eta G_\alpha}{(1 - \alpha)^2}, \end{aligned} \quad (103)$$

with an empty sum when $t = 1$. This proves (96).

Finally, $\hat{u}_t = u_t(\mathbf{m}_t)$ and the feasibility of fixed SDAC policies imply

$$\begin{aligned} 0 \leq u_t(\mathbf{m}_t) - u_t &= [u_t(\mathbf{m}_t) - \alpha x_t]_+ \\ &\leq [\alpha x_t(\mathbf{m}_t) - \alpha x_t]_+ \\ &\leq \alpha d_t. \end{aligned} \quad (104)$$

Combining (104) with (96) proves (97). $\square$

Lemma 12 (Stability term). *For all horizons T ,*

$$\sum_{t=1}^T (c_t(x_t, u_t) - \ell_t(\mathbf{m}_t)) \leq \frac{\alpha \eta T G_\alpha (L_x + \alpha L_u)}{(1 - \alpha)^2}.$$

Proof. Both state–action pairs lie in $\mathcal{D}_\alpha$. Hence (5) and Lemma 11 give

$$\begin{aligned} c_t(x_t, u_t) - \ell_t(\mathbf{m}_t) &\leq L_x |x_t - x_t(\mathbf{m}_t)| + L_u |u_t - u_t(\mathbf{m}_t)| \\ &\leq \frac{\alpha \eta G_\alpha L_x}{(1-\alpha)^2} + \frac{\alpha^2 \eta G_\alpha L_u}{(1-\alpha)^2} \\ &= \frac{\alpha \eta G_\alpha (L_x + \alpha L_u)}{(1-\alpha)^2}. \end{aligned} \quad (105)$$

Summing (105) over $t = 1, \dots, T$ proves the claim. $\square$

Proof of Theorem 1. Define the stability term

$$\text{Stab}_T \triangleq \sum_{t=1}^T (c_t(x_t, u_t) - \ell_t(\mathbf{m}_t)).$$

Using $\ell_t(\mathbf{m}) = c_t(x_t(\mathbf{m}), u_t(\mathbf{m}))$, the regret decomposes as

$$\text{Reg}_T^{\text{SDAC}} = \text{Stab}_T + \sum_{t=1}^T \ell_t(\mathbf{m}_t) - \min_{\mathbf{m} \in \mathcal{M}_H} \sum_{t=1}^T \ell_t(\mathbf{m}). \quad (106)$$

We first bound the initial Bregman divergence. Fix $\mathbf{m} \in \mathcal{M}_H$. Since $\mathbf{m}_1 = \frac{1}{H+1} \mathbf{1}$,

$$\begin{aligned} D_\Psi(\mathbf{m} \parallel \mathbf{m}_1) &= \sum_{i=1}^H m_i \log((H+1)m_i) \\ &\quad - \|\mathbf{m}\|_1 + \frac{H}{H+1} \\ &= \sum_{i=1}^H m_i \log m_i + \|\mathbf{m}\|_1 \log(H+1) \\ &\quad - \|\mathbf{m}\|_1 + \frac{H}{H+1}. \end{aligned} \quad (107)$$

Because $z \log z \leq 0$ for $z \in [0, 1]$,

$$D_\Psi(\mathbf{m} \parallel \mathbf{m}_1) \leq \|\mathbf{m}\|_1 \log(H+1) - \|\mathbf{m}\|_1 + \frac{H}{H+1}. \quad (108)$$

The right-hand side is affine in $\|\mathbf{m}\|_1 \in [0, 1]$, and its values at the endpoints are

$$\begin{aligned} \frac{H}{H+1} &\leq \log(H+1), \\ \log(H+1) - \frac{1}{H+1} &\leq \log(H+1). \end{aligned}$$

Here the first inequality follows from $\log z \geq 1 - z^{-1}$ for $z \geq 1$. Therefore,

$$D_\Psi(\mathbf{m} \parallel \mathbf{m}_1) \leq \log(H+1) \quad \text{for every } \mathbf{m} \in \mathcal{M}_H. \quad (109)$$

Lemma 9, $\|\mathbf{g}_t\|_\infty \leq G_\alpha$, and (109) imply, for every $\mathbf{m} \in \mathcal{M}_H$,

$$\begin{aligned} \sum_{t=1}^T (\ell_t(\mathbf{m}_t) - \ell_t(\mathbf{m})) &\leq \frac{D_\Psi(\mathbf{m} \parallel \mathbf{m}_1)}{\eta} + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{g}_t\|_\infty^2 \\ &\leq \frac{\log(H+1)}{\eta} + \frac{\eta G_\alpha^2 T}{2}. \end{aligned} \quad (110)$$

Since (110) holds uniformly over $\mathbf{m}$,

$$\sum_{t=1}^T \ell_t(\mathbf{m}_t) - \min_{\mathbf{m} \in \mathcal{M}_H} \sum_{t=1}^T \ell_t(\mathbf{m}) \leq \frac{\log(H+1)}{\eta} + \frac{\eta G_\alpha^2 T}{2}. \quad (111)$$

Lemma 12 gives

$$\text{Stab}_T \leq \frac{\alpha \eta T G_\alpha (L_x + \alpha L_u)}{(1-\alpha)^2}. \quad (112)$$

Substituting (111) and (112) into (106),

$$\begin{aligned} \text{Reg}_T^{\text{SDAC}} &\leq \frac{\log(H+1)}{\eta} + \frac{\eta T G_\alpha^2}{2} \\ &\quad + \frac{\alpha \eta T G_\alpha (L_x + \alpha L_u)}{(1-\alpha)^2} \\ &= \frac{\log(H+1)}{\eta} + \eta T \Gamma_\alpha. \end{aligned} \quad (113)$$

Differentiating with respect to η yields

$$-\frac{\log(H+1)}{\eta^2} + T \Gamma_\alpha.$$

The unique minimizer is therefore

$$\eta^* = \sqrt{\frac{\log(H+1)}{T \Gamma_\alpha}}. \quad (114)$$

At this value the two summands in (113) are equal, so

$$\text{Reg}_T^{\text{SDAC}} \leq 2\sqrt{\Gamma_\alpha T \log(H+1)}. \quad (115)$$

It remains to verify the dependence on α . As $\alpha \downarrow 0$,

$$\begin{aligned} \frac{G_\alpha}{\alpha} &= \frac{L_x}{1-\alpha} + L_u \longrightarrow L_x + L_u, \\ \frac{\Gamma_\alpha}{\alpha^2} &= \frac{1}{2} \left(\frac{G_\alpha}{\alpha} \right)^2 + \frac{G_\alpha}{\alpha} \frac{L_x + \alpha L_u}{(1-\alpha)^2} \\ &\longrightarrow \frac{(L_x + L_u)^2}{2} + L_x(L_x + L_u) > 0. \end{aligned}$$

Therefore $\sqrt{\Gamma_\alpha} = \Theta(\alpha)$ as $\alpha \downarrow 0$. As $\alpha \uparrow 1$, write $\varepsilon = 1 - \alpha$. Then

$$\begin{aligned} \varepsilon G_\alpha &= \alpha(L_x + \varepsilon L_u) \longrightarrow L_x, \\ \varepsilon^3 \Gamma_\alpha &= \frac{\varepsilon^3 G_\alpha^2}{2} + \alpha(\varepsilon G_\alpha)(L_x + \alpha L_u) \\ &\longrightarrow L_x(L_x + L_u) > 0. \end{aligned}$$

Thus $\sqrt{\Gamma_\alpha} = \Theta((1-\alpha)^{-3/2})$ as $\alpha \uparrow 1$. $\square$

REFERENCES

- [1] N. Agarwal, B. Bullins, E. Hazan, S. Kakade, and K. Singh, "Online control with adversarial disturbances," in *Proceedings of the 36th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 111–119.
- [2] N. Agarwal, E. Hazan, and K. Singh, "Logarithmic regret for online control," *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [3] E. Hazan and K. Singh, "Introduction to Online Control," Apr. 2026, arXiv:2211.09619, revised April 2026. [Online]. Available: <https://arxiv.org/abs/2211.09619>
- [4] M. Simchowitz, K. Singh, and E. Hazan, "Improper learning for non-stochastic control," in *Proceedings of the Thirty Third Conference on Learning Theory*, ser. Proceedings of Machine Learning Research, vol. 125. PMLR, 2020, pp. 3320–3436.

- [5] S. Lale, K. Azizzadenesheli, B. Hassibi, and A. Anandkumar, "Logarithmic regret bound in partially observable linear dynamical systems," *Advances in Neural Information Processing Systems*, vol. 33, pp. 20 876–20 888, 2020.
- [6] E. Hazan, S. M. Kakade, K. Singh, and A. Sidford, "Nonstochastic control with bandit feedback," *Advances in Neural Information Processing Systems*, vol. 33, pp. 16 774–16 785, 2020.
- [7] Y. Li, S. Das, and N. Li, "Online optimal control with affine constraints," *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 10, pp. 8527–8537, 2021. [Online]. Available: <https://ojs.aaai.org/index.php/AAAI/article/view/17035>
- [8] X. Liu, Z. Yang, and L. Ying, "Online nonstochastic control with adversarial and static constraints," in *Proceedings of the 40th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 22 277–22 288. [Online]. Available: <https://proceedings.mlr.press/v202/liu23at.html>
- [9] R. Kumar, S. Dean, and R. Kleinberg, "Online convex optimization with unbounded memory," 2024, arXiv:2210.09903, version 5. [Online]. Available: <https://arxiv.org/abs/2210.09903>
- [10] D. Shaviv and A. Özgür, "Universally near optimal online power control for energy harvesting nodes," *IEEE Journal on Selected Areas in Communications*, vol. 34, no. 12, pp. 3620–3631, 2016. [Online]. Available: <https://arxiv.org/abs/1511.00353>
- [11] A. Arafa, A. Baknina, and S. Ulukus, "Online fixed fraction policies in energy harvesting communication systems," *IEEE Transactions on Wireless Communications*, vol. 17, no. 5, pp. 2975–2986, 2018.
- [12] A. Zibaeenejad and J. Chen, "The optimal power control policy for an energy harvesting system with look-ahead: Bernoulli energy arrivals," in *2019 IEEE International Symposium on Information Theory (ISIT)*, 2019, extended version: arXiv:1904.12281. [Online]. Available: <https://arxiv.org/abs/1904.12281>
- [13] Q. Lin, J. Su, and M. Chen, "Optimal algorithms for online Age-of-Information optimization in energy harvesting systems," *IEEE Transactions on Networking*, vol. 33, no. 6, pp. 3146–3161, 2025.
- [14] H. Yu and M. J. Neely, "Learning-aided optimization for energy-harvesting devices with outdated state information," *IEEE/ACM Transactions on Networking*, vol. 27, no. 4, pp. 1501–1514, 2019.
- [15] K. Asgari and M. J. Neely, "Bregman-style online convex optimization with energy harvesting constraints," *Proceedings of the ACM on Measurement and Analysis of Computing Systems*, vol. 4, no. 3, Nov. 2020. [Online]. Available: <https://doi.org/10.1145/3428337>
- [16] L. Tassiulas and A. Ephremides, "Stability properties of constrained queueing systems and scheduling policies for maximum throughput in multihop radio networks," in *29th IEEE Conference on Decision and Control*. IEEE, 1990, pp. 2130–2132.
- [17] —, "Dynamic server allocation to parallel queues with randomly varying connectivity," *IEEE Transactions on Information Theory*, vol. 39, no. 2, pp. 466–478, 1993.
- [18] N. McKeown, A. Mekkittikul, V. Anantharam, and J. Walrand, "Achieving 100% throughput in an input-queued switch," *IEEE Transactions on Communications*, vol. 47, no. 8, pp. 1260–1267, 1999.
- [19] M. J. Neely, *Stochastic Network Optimization with Application to Communication and Queueing Systems*. Morgan & Claypool Publishers, 2010.
- [20] M. A. Khojastepour and A. Sabharwal, "Delay-constrained scheduling: Power efficiency, filter design, and bounds," in *IEEE INFOCOM 2004*, vol. 3. IEEE, 2004, pp. 1938–1949.
- [21] L. Yang, M. H. Hajiesmaili, R. Sitaraman, A. Wierman, E. Mallada, and W. S. Wong, "Online linear optimization with inventory management constraints," *Proceedings of the ACM on Measurement and Analysis of Computing Systems*, vol. 4, no. 1, 2020.
- [22] M. Hihat and A. Fermanian, "Online policy selection for inventory problems," *arXiv preprint arXiv:2411.19269*, 2024. [Online]. Available: <https://arxiv.org/abs/2411.19269>
- [23] P. A. P. Moran, "A probability theory of dams and storage systems," *Australian Journal of Applied Science*, vol. 5, no. 2, pp. 116–124, 1954.
- [24] N. U. Prabhu, *Stochastic Storage Processes: Queues, Insurance Risk, Dams, and Data Communication*. New York, NY: Springer New York, 1998. [Online]. Available: <https://doi.org/10.1007/978-1-4612-1742-8>
- [25] R. T. Rockafellar and R. J.-B. Wets, *Variational Analysis*, ser. Grundlehren der mathematischen Wissenschaften. Berlin: Springer, 1998, vol. 317.
- [26] E. Hazan, "Introduction to online convex optimization," *Foundations and Trends in Optimization*, vol. 2, no. 3–4, pp. 157–325, 2016.