

Online Constrained Control of Storage Systems via Simplex Disturbance-Action Policies

Kamiar Asgari and Michael J. Neely

Abstract—We study online control of a scalar storage system under nonnegative adversarial arrivals, convex state–action costs, and state-dependent action constraints. At each time, the controller chooses a feasible action before the current arrival and cost are revealed. We introduce *Simplex Disturbance-Action Control* (SDAC), a simplex-constrained finite-memory policy class whose fixed policies satisfy the per-step action constraint by construction. For online adaptation, an entropic online mirror descent (OMD) update maintains the SDAC parameters on the subprobability simplex, while scalar action clipping ensures feasibility of the actual time-varying trajectory. We prove $O(\sqrt{T \log(H+1)})$ regret against the best fixed SDAC policy, where H is the memory horizon. We further show that SDAC geometrically approximates every feasible proportional state-feedback policy; choosing $H = \Theta(\log T)$ yields $O(\sqrt{T \log \log T})$ regret against the best policy in this class. The setting is motivated by energy-harvesting battery management and related storage systems.

Index Terms—Online control, disturbance-action policies, storage systems, energy harvesting, regret minimization, state-dependent action constraints.

I. INTRODUCTION

A. Motivation and Problem Setting

MANY dynamic resource systems can be modeled by a storage state that evolves over time and constrains the control action at every time. The controller must choose an action subject to a state-dependent action constraint before observing the corresponding arrival and cost.

This creates a constrained discrete-time control problem with state–action coupling: aggressive early actions can harm later performance, while overly conservative actions can increase the stage cost. Our objective is to remain feasible at every time while competing with the best fixed policy in a tractable class that approximates feasible proportional state-feedback policies.

B. Model and Information Structure

We study a discrete-time storage system with state $x_t \geq 0$, where x_t models the storage level at the beginning of time t . At each time t , the controller chooses a control action u_t under the state-dependent action constraint

$$0 \leq u_t \leq \alpha x_t,$$

where $\alpha \in (0, 1)$ is a fixed known retention coefficient (so a fraction $1 - \alpha$ of the state is lost each period). After the action

is chosen, a nonnegative adversarial arrival w_t (a resource inflow) is revealed, and the state evolves according to

$$x_{t+1} = \alpha x_t - u_t + w_t.$$

Although this input is commonly called a disturbance in the online-control literature, we call $w_t \in [0, 1]$ an *arrival* to emphasize its nonnegativity; we retain the standard names Disturbance-Action Control (DAC) and Simplex Disturbance-Action Control (SDAC). Nonnegativity is essential to the storage interpretation and to the feasibility properties of SDAC.

Unlike stochastic models, our formulation is adversarial and non-anticipatory. At time t , the environment may select w_t and c_t using the history through time $t - 1$, but not the current action; the controller then chooses u_t , after which w_t and the entire function c_t are revealed. The goal is low cumulative cost relative to the best fixed policy in a class defined below, with both the online and benchmark trajectories remaining feasible.

C. Application Domains

This abstraction captures a broad set of domains:

- **Energy harvesting and battery management:** x_t is battery state of charge (retention coefficient α after self-discharge), w_t is harvested energy, and u_t is discharge or consumption bounded by available energy αx_t .
- **Queueing networks and buffer management:** x_t is queue/buffer occupancy, w_t is packet arrivals, and u_t is service bounded by αx_t .
- **Inventory control:** x_t is on-hand inventory (retention coefficient α after spoilage/deterioration), w_t is supply or replenishment, and u_t is sales bounded by αx_t .
- **Water-reservoir management:** x_t is stored water (retention coefficient α after evaporation/seepage), w_t is inflow, and u_t is controlled release bounded by αx_t .

We take energy harvesting as the running interpretation; the other domains share the same storage dynamics and state-dependent action constraint.

D. Contributions

Our contributions are as follows:

- **Policy class.** We introduce SDAC, a DAC-inspired policy class for storage dynamics with the state-dependent action constraint $0 \leq u_t \leq \alpha x_t$. We prove that every fixed SDAC policy is feasible, its induced loss is convex and Lipschitz in the policy parameter, and the class geometrically approximates feasible proportional state-feedback policies.

K. Asgari and M. J. Neely are with the Department of Electrical and Computer Engineering, University of Southern California, Los Angeles, CA 90089 USA (e-mail: kamiaras@usc.edu; mjneely@usc.edu).

Corresponding author: K. Asgari.

- **Online control and regret.** We develop an entropic OMD controller over SDAC parameters. A Bregman projection maintains the simplex constraint on the parameters, and scalar action clipping maintains feasibility of the actual trajectory. The controller achieves $O(\sqrt{T \log(H+1)})$ regret against the best fixed SDAC policy and, with $H = \Theta(\log T)$, achieves $O(\sqrt{T \log \log T})$ regret against the best feasible proportional state-feedback policy.

E. Related Work

a) *Online nonstochastic control:* A recent line of work studies online control of linear dynamical systems under adversarial disturbances and costs. Its performance criterion is cumulative regret relative to a benchmark policy class chosen in hindsight, under suitable stability conditions [1]–[3]; see [4] for an overview.

A central idea in this literature is *Disturbance-Action Control* (DAC),¹ a policy parameterization that combines a stabilizing state-feedback term with a finite linear filter of past disturbances. Under standard assumptions, this representation yields tractable surrogate losses in the policy parameters and supports efficient no-regret updates. Extensions include partial observability [5] and bandit feedback [6].

Constrained online-control methods have also been developed for affine and general static state–input constraints [7], [8]. Our focus is the storage-specific constraint $0 \leq u_t \leq \alpha x_t$, whose feasible action set depends on the current state. Exploiting nonnegative arrivals, SDAC uses a simplex parameterization to make every fixed policy feasible; scalar action clipping preserves feasibility when the parameter changes online.

b) *Energy harvesting:* Lin et al. [9] study online age-of-information optimization in an energy-harvesting system with related battery dynamics. They compare with an unrestricted offline optimum through competitive ratio; their objective depends on the energy allocation but not directly on the battery state, and the current input is revealed before allocation. In our model, the controller acts before observing the current arrival and cost, the cost may depend jointly on state and action, and performance is measured by regret against a fixed policy class.

Yu and Neely [10] study expected utility maximization under i.i.d. system states. Asgari and Neely [11] express the guarantee as an $O(\sqrt{T})$ regret–battery-capacity tradeoff and improve its dependence on the decision dimension, allowing arbitrary loss sequences while retaining i.i.d. energy arrivals. Both works use action-only objectives and fixed-decision comparators. Here, arrivals are adversarial, costs depend on both state and action, and the comparator is a fixed state-responsive policy.

c) *Queueing and network optimization:* Queueing control provides a mathematically related perspective based on stochastic arrivals and real-time service decisions. Foundational work emphasizes throughput and queue stability [12]–[14], while stochastic network optimization also studies tradeoffs between time-average objectives and backlog [15]. Our

setting differs in its adversarial arrivals, convex state–action costs, and regret benchmark against a policy chosen in hindsight.

In many queueing applications, the natural retention coefficient is $\alpha = 1$ (no abandonment or leakage), so those systems are not direct instances of the $\alpha \in (0, 1)$ model analyzed here. Section VII discusses the additional issues that arise when $\alpha \geq 1$.

Among other results, Khojastepour and Sabharwal [16] study a time-invariant robust scheduler representable as a linear filter of past arrivals. It is designed to minimize average transmit power subject to a maximum packet-delay guarantee, rather than updated online to minimize revealed general costs.

d) *Inventory management:* Yang et al. [17] study online procurement under inventory constraints. The controller observes the current demand and price, then procures supply while ensuring that demand is met; performance is measured by competitive ratio for aggregate procurement cost. Our controller instead chooses an outflow before observing the current inflow and cost, and permits convex state–action costs.

Hihat and Fermanian [18] consider an inventory-control formulation with potentially nonlinear and nondifferentiable state-transition functions and study base-stock policies. While that model is more general in scope, its decision structure differs from ours: as in [17], the controller chooses supply and then observes demand through the transition, whereas our controller chooses an outflow and then observes incoming supply. Our scalar linear recursion and simplex policy class yield explicit feasibility and regret guarantees; the two frameworks target different dynamics, decisions, and policy classes.

e) *Water-reservoir management:* Classical dam theory studies storage dynamics with random inflows and controlled releases (see, for example, [19], [20]). Relative to our setting, this literature assumes stochastic inputs and emphasizes stationary or long-run performance quantities. It provides historical context for storage control under state constraints.

F. Paper Organization

Section III formalizes the adversarial storage-control model, assumptions, and regret benchmark. Section IV introduces the SDAC policy class and summarizes its structural properties. Section V presents the online controller, Section VI develops regret guarantees, Section VII discusses possible extensions, and Section VIII concludes. Deferred proofs appear in the Appendix.

II. NOTATION

We write $[z]_+ \triangleq \max\{z, 0\}$. For $v \in \mathbb{R}^n$, the i -th coordinate is $v^{(i)}$. We write $\mathbf{1}$ for the all-ones vector (dimension is clear from context). The Euclidean inner product is denoted by $\langle \cdot, \cdot \rangle$. We set $w_\tau \equiv 0$ for $\tau \leq 0$. We write $\mathcal{A}$ for the online algorithm and π for a fixed policy (also written $\pi(M)$ when parameterized by a vector M).

¹“Disturbance-Action Control” is the standard name in the online-control literature. In this paper the underlying signal is the nonnegative arrival w_t ; we retain the DAC/SDAC names for continuity with that literature.

III. MODEL AND PROBLEM SETUP

We consider a discrete-time process indexed by $t \in \{1, \dots, T\}$. The system state $x_t \in \mathbb{R}_{\geq 0}$ denotes the resource level at the beginning of time t . The initial state is given as

$$x_1 = 0. \quad (1)$$

Fix a retention coefficient $\alpha \in (0, 1)$. At each time t , the controller chooses a control action $u_t \in \mathbb{R}_{\geq 0}$ subject to the state-dependent action constraint

$$0 \leq u_t \leq \alpha x_t. \quad (2)$$

The system then receives an arrival $w_t \in [0, 1]$ and evolves according to

$$x_{t+1} = \alpha x_t - u_t + w_t. \quad (3)$$

For every feasible trajectory,

$$0 \leq x_{t+1} \leq \alpha x_t + 1.$$

Since $x_1 = 0$, induction gives

$$0 \leq x_t \leq \frac{1 - \alpha^{t-1}}{1 - \alpha} \leq \frac{1}{1 - \alpha}, \quad 0 \leq u_t \leq \frac{\alpha}{1 - \alpha}.$$

Define the reachable state-action domain

$$\mathcal{D}_\alpha \triangleq \left\{ (x, u) : 0 \leq x \leq \frac{1}{1 - \alpha}, 0 \leq u \leq \alpha x \right\}. \quad (4)$$

After the control action is selected, a convex cost function

$$c_t : \mathbb{R}_{\geq 0}^2 \rightarrow \mathbb{R}$$

is revealed, and the incurred cost is $c_t(x_t, u_t)$. We assume each c_t is convex in (x, u) and satisfies the following Lipschitz condition on $\mathcal{D}_\alpha$: there exist constants $L_x, L_u > 0$, independent of t , such that for all $(x, u), (x', u') \in \mathcal{D}_\alpha$,

$$|c_t(x, u) - c_t(x', u')| \leq L_x |x - x'| + L_u |u - u'|. \quad (5)$$

We also assume access, at every $(x, u) \in \mathcal{D}_\alpha$, to a cost subgradient $s_t(x, u) = (s_{x,t}, s_{u,t}) \in \partial c_t(x, u)$ satisfying $|s_{x,t}| \leq L_x$ and $|s_{u,t}| \leq L_u$. This oracle condition holds, for example, when the Lipschitz property above extends to an open neighborhood of $\mathcal{D}_\alpha$.

The information pattern is adversarial and non-anticipatory. At time t , the environment may choose (w_t, c_t) as a function of the history through time $t - 1$, but not of the current action (or current randomization). The controller chooses u_t without observing (w_t, c_t) ; the pair is then revealed. Subject to this timing and the stated regularity conditions, the arrival and cost sequences may be arbitrary.

a) Assumptions:

- 1) **Retention coefficient:** $\alpha \in (0, 1)$ is fixed and known (a fraction $1 - \alpha$ of the state is lost each period).
- 2) **Arrival bounds:** $w_t \in [0, 1]$ for all $t \in \{1, \dots, T\}$. In particular, arrivals are nonnegative resource increments, in contrast to signed process-noise disturbances common in classical control.
- 3) **Adversarial timing:** (w_t, c_t) may depend on the history through time $t - 1$ but not on the current action, and is revealed only after u_t is chosen.

- 4) **Cost regularity:** each c_t is convex in (x, u) , obeys (5) on $\mathcal{D}_\alpha$, and admits the bounded subgradient oracle described above.
- 5) **Feasible policy class:** both the online policy and the benchmark policies are required to satisfy (2) under (3).

We focus on $\alpha \in (0, 1)$; the case $\alpha \geq 1$ is discussed in Section VII.

Let $\{(x_t^A, u_t^A)\}_{t=1}^T$ denote the trajectory of an online algorithm $\mathcal{A}$. Define its cumulative cost by

$$J_T(\mathcal{A}) \triangleq \sum_{t=1}^T c_t(x_t^A, u_t^A). \quad (6)$$

For a fixed positive integer H , let π range over the policy class SDAC_H (defined in Section IV). For each $\pi \in \text{SDAC}_H$, let $\{(x_t^\pi, u_t^\pi)\}_{t=1}^T$ be the induced feasible trajectory under the same arrival sequence. The regret against the best fixed SDAC policy is defined by

$$\text{Reg}_{\text{SDAC}}(T) \triangleq J_T(\mathcal{A}) - \min_{\pi \in \text{SDAC}_H} \sum_{t=1}^T c_t(x_t^\pi, u_t^\pi). \quad (7)$$

While the online controller may update its SDAC parameters over time, the benchmark in (7) is the best single fixed SDAC policy chosen in hindsight. The objective is to design an online policy that satisfies (2) for all $t \in \{1, \dots, T\}$ and achieves sublinear regret as defined in (7).

IV. SIMPLEX DISTURBANCE-ACTION CONTROL (SDAC)

A. From Disturbance-Action Control to SDAC

Disturbance-Action Control (DAC) combines a stabilizing state-feedback term with a finite linear filter of past disturbances. Under standard conditions, this parameterization yields convex surrogate losses and efficient online updates. In our scalar storage model, specializing the baseline feedback term to zero gives the representative form

$$u_t = \sum_{i=1}^H M^{(i)} w_{t-i}, \quad (8)$$

where $H \in \mathbb{N}$ is the memory horizon and $M \in \mathbb{R}^H$ is a linear filter. We allow nonpositive indices in such sums, with the convention $w_\tau \equiv 0$ for $\tau \leq 0$.

For our constrained model, an unrestricted filter is insufficient: its action can violate (2) by exceeding αx_t or becoming negative. We therefore introduce *Simplex Disturbance-Action Control* (SDAC), which uses a simplex-constrained filter tailored to the storage dynamics. This restriction yields fixed-policy feasibility while preserving convex surrogate losses and tractable online updates.

B. SDAC Definition

Fix a memory horizon $H \in \mathbb{N}$. Let

$$\Delta_H \triangleq \left\{ M \in \mathbb{R}_{\geq 0}^H : \sum_{i=1}^H M^{(i)} \leq 1 \right\} \quad (9)$$

be the subprobability simplex. For any $M \in \Delta_H$, define the policy $\pi(M)$ by

$$u_t^{\pi(M)} \triangleq \sum_{i=1}^H M^{(i)} \alpha^i w_{t-i}, \quad w_\tau \equiv 0 \text{ for } \tau \leq 0. \quad (10)$$

For each time $t \in \{1, \dots, T\}$, let

$$\tilde{w}_t \triangleq (\alpha w_{t-1}, \alpha^2 w_{t-2}, \dots, \alpha^H w_{t-H})^\top \in \mathbb{R}^H, \quad (11)$$

with the same convention $w_\tau \equiv 0$ for $\tau \leq 0$, so that (10) can be written as

$$u_t^{\pi(M)} = \langle M, \tilde{w}_t \rangle. \quad (12)$$

The induced state trajectory satisfies

$$x_{t+1}^{\pi(M)} = \alpha x_t^{\pi(M)} - u_t^{\pi(M)} + w_t, \quad x_1^{\pi(M)} = 0. \quad (13)$$

The weights α^i align the arrival filter with the retention coefficient α in (3).

We now establish the two structural properties that make SDAC useful: the induced control action always satisfies the feasibility constraint (2), and the induced cost is convex and Lipschitz in M , so that regret minimization against the best fixed SDAC policy is an instance of online convex optimization over the simplex. Both properties follow from an explicit decomposition of the state trajectory that isolates its dependence on M .

Recall that we set $w_\tau \equiv 0$ for $\tau \leq 0$, which is used in the base-case arguments below.

Lemma 1 (State decomposition under SDAC). *Fix $H \in \mathbb{N}$ and $M \in \Delta_H$, and write*

$$W \triangleq \sum_{i=1}^H M^{(i)} \leq 1.$$

For $i \in \{1, \dots, H\}$, define

$$\psi_t^{(i)} \triangleq \sum_{j=1}^i \alpha^{j-1} w_{t-j}, \quad \rho_t \triangleq \sum_{j=1}^{t-1} \alpha^{j-1} w_{t-j},$$

where the second sum is empty when $t = 1$. Then, for every $t \in \{1, \dots, T\}$,

$$x_t^{\pi(M)} = \sum_{i=1}^H M^{(i)} \psi_t^{(i)} + (1 - W) \rho_t. \quad (14)$$

Moreover,

$$0 \leq \psi_t^{(i)} \leq \rho_t \leq \frac{1}{1 - \alpha},$$

and $M \mapsto x_t^{\pi(M)}$ is affine.

Proof. The auxiliary quantities satisfy

$$\rho_{t+1} = \alpha \rho_t + w_t, \quad \rho_1 = 0, \quad (15)$$

and, for every $i \in \{1, \dots, H\}$,

$$\begin{aligned} \alpha \psi_t^{(i)} - \alpha^i w_{t-i} &= \sum_{j=1}^{i-1} \alpha^j w_{t-j} \\ &= \psi_{t+1}^{(i)} - w_t. \end{aligned} \quad (16)$$

For $i = 1$, the sum in the first line is empty.

At $t = 1$, one has $x_1^{\pi(M)} = \psi_1^{(i)} = \rho_1 = 0$, so (14) holds. Suppose it holds at time t . Using (13) and (10),

$$\begin{aligned} x_{t+1}^{\pi(M)} &= \alpha \left(\sum_{i=1}^H M^{(i)} \psi_t^{(i)} + (1 - W) \rho_t \right) \\ &\quad - \sum_{i=1}^H M^{(i)} \alpha^i w_{t-i} + w_t \\ &= \sum_{i=1}^H M^{(i)} (\alpha \psi_t^{(i)} - \alpha^i w_{t-i}) \\ &\quad + \alpha(1 - W) \rho_t + w_t. \end{aligned} \quad (17)$$

Substituting (15) and (16) into (17),

$$\begin{aligned} x_{t+1}^{\pi(M)} &= \sum_{i=1}^H M^{(i)} (\psi_{t+1}^{(i)} - w_t) \\ &\quad + (1 - W)(\rho_{t+1} - w_t) + w_t \\ &= \sum_{i=1}^H M^{(i)} \psi_{t+1}^{(i)} + (1 - W) \rho_{t+1}, \end{aligned} \quad (18)$$

because the total coefficient of w_t in the first line is $-W - (1 - W) + 1 = 0$. This proves (14) by induction.

All summands defining $\psi_t^{(i)}$ and ρ_t are nonnegative. Since $w_\tau = 0$ for $\tau \leq 0$, $\psi_t^{(i)}$ is a truncation of ρ_t ; hence

$$0 \leq \psi_t^{(i)} \leq \rho_t \leq \sum_{j=1}^{\infty} \alpha^{j-1} = \frac{1}{1 - \alpha}.$$

Finally,

$$x_t^{\pi(M)} = \rho_t + \sum_{i=1}^H M^{(i)} (\psi_t^{(i)} - \rho_t),$$

which is affine in M . $\square$

The first consequence of this decomposition is fixed-policy feasibility.

Lemma 2 (Feasibility of SDAC policies). *For every $M \in \Delta_H$, the trajectory induced by (10) and (13) satisfies*

$$0 \leq u_t^{\pi(M)} \leq \alpha x_t^{\pi(M)} \quad \text{for all } t \in \{1, \dots, T\}. \quad (19)$$

Proof. Because $M^{(i)}$, α , and w_{t-i} are nonnegative,

$$u_t^{\pi(M)} = \sum_{i=1}^H M^{(i)} \alpha^i w_{t-i} \geq 0.$$

For every $i \in \{1, \dots, H\}$,

$$\alpha \psi_t^{(i)} = \sum_{j=1}^i \alpha^j w_{t-j} \geq \alpha^i w_{t-i}. \quad (20)$$

Let $W = \sum_{i=1}^H M^{(i)}$. By Lemma 1, $1 - W \geq 0$ and $\rho_t \geq 0$. Therefore,

$$\begin{aligned} \alpha x_t^{\pi(M)} &= \sum_{i=1}^H M^{(i)} \alpha \psi_t^{(i)} + \alpha(1 - W) \rho_t \\ &\geq \sum_{i=1}^H M^{(i)} \alpha \psi_t^{(i)} \\ &\geq \sum_{i=1}^H M^{(i)} \alpha^i w_{t-i} = u_t^{\pi(M)}, \end{aligned} \quad (21)$$

where the second inequality uses (20). $\square$

For fixed arrivals, both coordinates of the SDAC state-action pair are affine in M . The following lemma records the resulting convexity and the constants used by the online update.

Lemma 3 (Convexity and Lipschitz continuity of surrogate cost). *Define*

$$\ell_t(M) \triangleq c_t(x_t^{\pi(M)}, u_t^{\pi(M)}), \quad M \in \Delta_H,$$

and

$$L \triangleq \alpha L_u + \frac{L_x}{1 - \alpha}. \quad (22)$$

Then ℓ_t is convex on Δ_H , and for all $M, M' \in \Delta_H$,

$$|\ell_t(M) - \ell_t(M')| \leq L \|M - M'\|_1.$$

Moreover, at every $M \in \Delta_H$, one can select $g \in \partial \ell_t(M)$ such that $\|g\|_\infty \leq L$.

Proof. For $i \in \{1, \dots, H\}$, define

$$a_t^{(i)} \triangleq \psi_t^{(i)} - \rho_t, \quad b_t^{(i)} \triangleq \alpha^i w_{t-i}.$$

Lemma 1 and (10) give

$$x_t^{\pi(M)} = \rho_t + \langle a_t, M \rangle, \quad u_t^{\pi(M)} = \langle b_t, M \rangle. \quad (23)$$

The coefficient bounds in Lemma 1, together with $w_{t-i} \in [0, 1]$, imply

$$\|a_t\|_\infty \leq \frac{1}{1 - \alpha}, \quad \|b_t\|_\infty \leq \alpha. \quad (24)$$

Since (23) is affine in M and c_t is convex, ℓ_t is convex.

Let $\Delta M \triangleq M - M'$. By (5) and (23),

$$\begin{aligned} |\ell_t(M) - \ell_t(M')| &\leq L_x |\langle a_t, \Delta M \rangle| + L_u |\langle b_t, \Delta M \rangle| \\ &\leq (L_x \|a_t\|_\infty + L_u \|b_t\|_\infty) \|\Delta M\|_1 \\ &\leq \left(\frac{L_x}{1 - \alpha} + \alpha L_u \right) \|\Delta M\|_1 \\ &= L \|M - M'\|_1. \end{aligned} \quad (25)$$

Choose $(s_x, s_u) \in \partial c_t(x_t^{\pi(M)}, u_t^{\pi(M)})$ with $|s_x| \leq L_x$ and $|s_u| \leq L_u$, and set

$$g \triangleq s_x a_t + s_u b_t. \quad (26)$$

For every $M' \in \Delta_H$, the subgradient inequality and (23) yield

$$\begin{aligned} \ell_t(M') &\geq \ell_t(M) + s_x \langle a_t, M' - M \rangle + s_u \langle b_t, M' - M \rangle \\ &= \ell_t(M) + \langle g, M' - M \rangle. \end{aligned}$$

Thus $g \in \partial \ell_t(M)$. Finally,

$$\begin{aligned} \|g\|_\infty &\leq |s_x| \|a_t\|_\infty + |s_u| \|b_t\|_\infty \\ &\leq \frac{L_x}{1 - \alpha} + \alpha L_u = L. \end{aligned} \quad (27)$$

$\square$

Regret minimization against the best fixed SDAC policy thus becomes online convex optimization over the simplex. optimization over the simplex.

C. Expressiveness: Proportional State-Feedback Approximation

A standard policy class in online control is the family of linear state-feedback policies $u_t = -Kx_t$, and a foundational property of Disturbance-Action Control is that the DAC class *approximates* any strongly stable linear policy, with approximation error decaying geometrically in the memory horizon H [1], [4].

Under our sign convention, positive u_t drains storage, so proportional feedback is written $u_t = qx_t$. In our scalar model, the feasible proportional state-feedback policies are exactly those with gain $q \in [0, \alpha]$: since $x_t \geq 0$, the feasibility constraint $0 \leq u_t \leq \alpha x_t$ in (2) holds if and only if $q \in [0, \alpha]$. The next lemma gives an explicit approximation within the feasible class SDAC_H ; the two endpoint policies are represented exactly.

Lemma 4 (SDAC approximation of proportional state feedback). *Fix $q \in [0, \alpha]$, set $\beta \triangleq \alpha - q$, and let π_q^{lin} denote the feasible linear policy*

$$\begin{aligned} u_t^{\text{lin}} &= q x_t^{\text{lin}}, \quad x_{t+1}^{\text{lin}} = \beta x_t^{\text{lin}} + w_t, \\ t &\in \{1, \dots, T\}, \quad x_1^{\text{lin}} = 0. \end{aligned}$$

Define $M_q^* \in \mathbb{R}^H$ by

$$M_q^{*(i)} \triangleq \frac{q \beta^{i-1}}{\alpha^i}, \quad i = 1, \dots, H.$$

Then $M_q^* \in \Delta_H$, and if each c_t satisfies (5), then for every $\{w_t\} \subset [0, 1]$,

$$\begin{aligned} \left| \sum_{t=1}^T c_t(x_t^{\pi(M_q^*)}, u_t^{\pi(M_q^*)}) - \sum_{t=1}^T c_t(x_t^{\text{lin}}, u_t^{\text{lin}}) \right| \\ \leq \left(L_u + \frac{L_x}{1 - \alpha} \right) \frac{q \beta^H}{1 - \beta} T. \end{aligned} \quad (28)$$

The proof is deferred to Appendix A.

Theorem 1 depends on the memory horizon only through $\sqrt{\log(H+1)}$. Combining that bound with Lemma 4 therefore yields a no-regret guarantee against the best feasible proportional state-feedback policy; Corollary 1 makes this statement precise.

V. ONLINE ALGORITHM

We now present our online algorithm, which updates SDAC parameters by online mirror descent (OMD) with unnormalized negentropy (equivalently, the exponentiated gradient method) and a fixed step size $\eta > 0$ over the subprobability simplex $\Delta_H \subset \mathbb{R}^H$ [21].

For each time $t \in \{1, \dots, T\}$, let $M_t \in \Delta_H$ be the online parameter. We initialize at

$$M_1^{(i)} \triangleq \frac{1}{H+1}, \quad i = 1, \dots, H \quad (29)$$

so that the maximum initial divergence $\sup_{M \in \Delta_H} D_\Psi(M \| M_1)$ is at most $\log(H+1)$ (see the proof of Theorem 1).

a) *Control-action selection:* The controller first forms the nominal SDAC action

$$\hat{u}_t \triangleq \langle M_t, \tilde{w}_t \rangle, \quad (30)$$

and clips it to the admissible interval $[0, \alpha x_t^A]$:

$$u_t^A \triangleq \min\{\hat{u}_t, \alpha x_t^A\}. \quad (31)$$

The state update is

$$x_{t+1}^A = \alpha x_t^A - u_t^A + w_t, \quad x_1^A = 0. \quad (32)$$

By construction, $\hat{u}_t \geq 0$ and $u_t^A \in [0, \alpha x_t^A]$; with $x_1^A = 0$ and $w_t \in [0, 1]$, induction on (32) yields $x_t^A \geq 0$ for all $t \in \{1, \dots, T\}$, so the online trajectory satisfies (2).

b) *Surrogate cost and OMD update:* After (w_t, c_t) is revealed, define the surrogate cost by

$$\ell_t(M) \triangleq c_t(x_t^{\pi(M)}, u_t^{\pi(M)}), \quad M \in \Delta_H, \quad (33)$$

which is convex and Lipschitz by Lemma 3. Select a subgradient $g_t \in \partial \ell_t(M_t)$ satisfying $\|g_t\|_\infty \leq L$, whose existence and construction are given in that lemma. Let

$$\Psi(M) \triangleq \sum_{i=1}^H (M^{(i)} \log M^{(i)} - M^{(i)}), \quad M \in \mathbb{R}_{\geq 0}^H, \quad (34)$$

denote the unnormalized negentropy (with the convention $0 \log 0 \triangleq 0$), and let

$$D_\Psi(M \| V) \triangleq \sum_{i=1}^H \left(M^{(i)} \log \frac{M^{(i)}}{V^{(i)}} - M^{(i)} + V^{(i)} \right)$$

be its Bregman divergence for $V \in \mathbb{R}_{>0}^H$, the generalized Kullback–Leibler divergence. Terms with $M^{(i)} = 0$ are interpreted by continuity as zero in the logarithmic part.

We define the OMD update

$$M_{t+1} \triangleq \arg \min_{M \in \Delta_H} \left\{ \eta \langle g_t, M \rangle + D_\Psi(M \| M_t) \right\}. \quad (35)$$

The regret analysis in Section VI works directly with this optimization-based definition.

c) *Closed-form solution:* Although (35) is a constrained optimization problem, it is easy to solve analytically because Ψ is separable and the objective is linear in g_t . As shown in Lemma 6, the minimizer in (35) is exactly the result of two elementary steps: a multiplicative step,

$$\tilde{M}_{t+1}^{(i)} \triangleq M_t^{(i)} \exp(-\eta g_t^{(i)}), \quad i = 1, \dots, H, \quad (36)$$

followed by the Bregman projection onto Δ_H given by (37) (the scaling is nontrivial only when $S_{t+1} > 1$):

$$M_{t+1}^{(i)} \triangleq \frac{\tilde{M}_{t+1}^{(i)}}{\max\{1, S_{t+1}\}}, \quad S_{t+1} \triangleq \sum_{j=1}^H \tilde{M}_{t+1}^{(j)}, \quad (37)$$

for $i = 1, \dots, H$. This is the form we implement in practice.

VI. REGRET ANALYSIS

The following theorem states that the online controller achieves sublinear regret against every fixed SDAC policy $\pi(M)$, where M is any parameter in Δ_H . The regret proof has two parts: a standard online convex optimization bound for the surrogate costs ℓ_t , and a stability bound showing that the action-clipped online trajectory stays close to the SDAC policy trajectory driven by the same parameters. Some technical lemmas required for the proof are deferred to Appendix A.

Theorem 1 (Regret bound). *Suppose the cost functions satisfy the regularity and bounded-subgradient assumptions of Section III. Then the OMD controller with action clipping from Section V, initialized as in (29) and using step size $\eta > 0$, satisfies*

$$\text{Reg}_{\text{SDAC}}(T) \leq \frac{\log(H+1)}{\eta} + \eta L^2 T \left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right),$$

where $L \triangleq \alpha L_u + \frac{L_x}{1-\alpha}$ as in (22). In particular, the choice

$$\eta^* = \sqrt{\frac{\log(H+1)}{L^2 T \left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right)}}$$

gives

$$\text{Reg}_{\text{SDAC}}(T) \leq 2L \sqrt{\left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right) T \log(H+1)}.$$

Combining Theorem 1 with Lemma 4 yields a uniform guarantee against the best feasible proportional state-feedback policy.

Corollary 1 (No-regret against proportional state feedback). *Set*

$$H_T \triangleq \max \left\{ 1, \left\lceil \frac{\log T}{\log(1/\alpha)} \right\rceil \right\}.$$

Run the OMD controller with action clipping using memory horizon $H = H_T$ and the step size η^ of Theorem 1. For each $q \in [0, \alpha]$, let $(x_t^{(q)}, u_t^{(q)})$ denote the trajectory of the proportional policy $u_t^{(q)} = q x_t^{(q)}$. Define*

$$\text{Reg}_{\text{lin}}(T) \triangleq \sum_{t=1}^T c_t(x_t^A, u_t^A) - \min_{q \in [0, \alpha]} \sum_{t=1}^T c_t(x_t^{(q)}, u_t^{(q)}).$$

Then

$$\begin{aligned} \text{Reg}_{\text{lin}}(T) &\leq 2L \sqrt{\left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right) T \log(H_T + 1)} \\ &\quad + \left(L_u + \frac{L_x}{1-\alpha} \right) \frac{\alpha^{H_T+1}}{1-\alpha} T, \end{aligned}$$

where $L = \alpha L_u + \frac{L_x}{1-\alpha}$. Since $\alpha^{H_T} \leq 1/T$, the second term is $O(1)$, while $\log(H_T + 1) = O(\log \log T)$ as $T \rightarrow \infty$. Consequently,

$$\text{Reg}_{\text{lin}}(T) = O\left(\sqrt{T \log \log T}\right).$$

Thus $\text{Reg}_{\text{lin}}(T) = o(T)$ against the best proportional state-feedback policy with $q \in [0, \alpha]$.

Proof. For $M \in \Delta_{H_T}$ and $q \in [0, \alpha]$, write

$$J_T(M) \triangleq \sum_{t=1}^T c_t(x_t^{\pi(M)}, u_t^{\pi(M)}),$$

$$J_T(q) \triangleq \sum_{t=1}^T c_t(x_t^{(q)}, u_t^{(q)}).$$

Let

$$R_T \triangleq 2L\sqrt{\left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2}\right) T \log(H_T + 1)}.$$

Fix $q \in [0, \alpha]$, set $\beta = \alpha - q$, and use the comparator M_q^* from Lemma 4. Then

$$\begin{aligned} J_T(\mathcal{A}) - J_T(q) &= \left(J_T(\mathcal{A}) - \min_{M \in \Delta_{H_T}} J_T(M) \right) \\ &\quad + \left(\min_{M \in \Delta_{H_T}} J_T(M) - J_T(q) \right) \\ &\leq R_T + J_T(M_q^*) - J_T(q) \\ &\leq R_T + \left(L_u + \frac{L_x}{1-\alpha} \right) \frac{q\beta^{H_T}}{1-\beta} T. \end{aligned} \quad (38)$$

The third line uses Theorem 1; the fourth uses Lemma 4. Since $q \leq \alpha$, $\beta \leq \alpha$, and $1 - \beta \geq 1 - \alpha$,

$$\frac{q\beta^{H_T}}{1-\beta} \leq \frac{\alpha^{H_T+1}}{1-\alpha}. \quad (39)$$

The right-hand side is independent of q . Taking the minimum comparator cost over $q \in [0, \alpha]$ in (38) and applying (39) proves the finite-time bound.

Finally,

$$H_T \geq \frac{\log T}{\log(1/\alpha)} \implies \alpha^{H_T} \leq \frac{1}{T}.$$

Also $H_T = \Theta(\log T)$, so $\log(H_T + 1) = O(\log \log T)$ as $T \rightarrow \infty$. These two estimates give the stated asymptotic rate. $\square$

VII. EXTENSIONS

We briefly discuss three possible extensions. Establishing their full regret bounds is left for future work.

a) Retention $\alpha \geq 1$. For $\alpha \geq 1$, choose a fixed baseline drain Kx_t such that $\bar{\alpha} \triangleq \alpha - K \in (0, 1)$, and write $u_t = Kx_t + v_t$. The residual dynamics become $x_{t+1} = \bar{\alpha}x_t - v_t + w_t$. An SDAC filter for v_t should therefore use the powers $\bar{\alpha}^i$, not α^i . If $0 \leq v_t \leq \bar{\alpha}x_t$, then $0 \leq u_t \leq (K + \bar{\alpha})x_t = \alpha x_t$. This reduction suggests an extension of the present analysis with transformed cost and Lipschitz constants; we do not pursue the details here.

b) Vector action / multi-task allocation. Consider a single storage state x_t with a vector action $u_t \in \mathbb{R}_{\geq 0}^n$ that allocates resources across n tasks, subject to $\mathbf{1}^\top u_t \leq \alpha x_t$, so that the total drain from storage is the sum of the task allocations. The dynamics remain $x_{t+1} = \alpha x_t - \mathbf{1}^\top u_t + w_t$, and the cost $c_t(x_t, u_t)$ is convex in (x, u) . The fixed-policy SDAC construction lifts by replacing the filter vector $M \in \Delta_H$ with a matrix $M \in \mathbb{R}_{\geq 0}^{n \times H}$ whose entries sum to at most one (a matrix simplex), with each row corresponding to a task and each column to a memory lag. OMD can then run over this matrix simplex, with radial action clipping when the nominal total allocation exceeds αx_t .

c) Multiple storages with diagonal dynamics. Suppose the state is $x_t \in \mathbb{R}_{\geq 0}^d$ and evolves as $x_{t+1} = Ax_t - u_t + w_t$ with $A = \text{diag}(\alpha_1, \dots, \alpha_d)$, $\alpha_i \in (0, 1)$, arrivals $w_t \in [0, 1]^d$, and per-coordinate constraints $0 \leq u_t^{(i)} \leq \alpha_i x_t^{(i)}$. The dynamics and feasible set then decouple by coordinate, and the natural SDAC class is a product of simplices. For a jointly convex cost, OMD operates jointly on this product set; the optimization separates by coordinate only when the cost is coordinate-separable. Cross-coupled storage systems (non-diagonal A) require additional analysis.

These constructions retain the basic simplex/OMD structure, but their complete feasibility and regret analyses require separate treatment.

VIII. CONCLUSION

We studied adversarial online control of a storage system under a state-dependent action constraint. Every fixed SDAC policy is feasible by construction. For the time-varying online controller, a Bregman projection maintains the SDAC parameter constraint, while scalar action clipping maintains feasibility of the actual trajectory. The resulting controller achieves $O(\sqrt{T \log(H+1)})$ regret against the best fixed SDAC policy. The SDAC approximation result further gives $O(\sqrt{T \log \log T})$ regret against the best feasible proportional state-feedback policy when $H = \Theta(\log T)$. The model applies directly to energy-harvesting battery management and related settings in which storage limits must be respected in real time. Section VII discusses possible extensions to $\alpha \geq 1$, vector allocation, and diagonal multi-storage systems.

APPENDIX A DEFERRED PROOFS

This appendix collects the deferred proof of the proportional-state-feedback approximation guarantee from Section IV, together with the OMD and stability lemmas leading up to the regret bound stated in Section VI.

Recall that we set $w_\tau \equiv 0$ for $\tau \leq 0$, which is used in several base-case arguments below.

We first prove the linear-policy approximation guarantee stated in Section IV-C.

Proof of Lemma 4. Fix $q \in [0, \alpha]$ and set $\beta \triangleq \alpha - q \in [0, \alpha]$.

Proportional-policy trajectory. Substituting $u_t^{\text{lin}} = qx_t^{\text{lin}}$ into the dynamics gives

$$x_{t+1}^{\text{lin}} = \beta x_t^{\text{lin}} + w_t, \quad x_1^{\text{lin}} = 0. \quad (40)$$

Unrolling (40),

$$\begin{aligned} x_t^{\text{lin}} &= \sum_{j=1}^{t-1} \beta^{j-1} w_{t-j}, \\ u_t^{\text{lin}} &= q \sum_{j=1}^{t-1} \beta^{j-1} w_{t-j}. \end{aligned} \quad (41)$$

Thus $x_t^{\text{lin}} \geq 0$, and $q \in [0, \alpha]$ implies

$$0 \leq u_t^{\text{lin}} = qx_t^{\text{lin}} \leq \alpha x_t^{\text{lin}}.$$

Simplex membership. Each coordinate of M_q^* is nonnegative. If $q = 0$, then $M_q^* = 0 \in \Delta_H$. Suppose $q > 0$ and define

$$r \triangleq \frac{\beta}{\alpha} \in [0, 1).$$

Since $q = \alpha - \beta$, one has $q/\alpha = 1 - r$. Consequently,

$$\begin{aligned} \sum_{i=1}^H M_q^{*(i)} &= \frac{q}{\alpha} \sum_{i=1}^H r^{i-1} \\ &= (1-r) \frac{1-r^H}{1-r} \\ &= 1-r^H \leq 1. \end{aligned} \quad (42)$$

Hence $M_q^* \in \Delta_H$.

Action approximation. The identity $M_q^{*(i)} \alpha^i = q \beta^{i-1}$ gives

$$u_t^{\pi(M_q^*)} = q \sum_{i=1}^H \beta^{i-1} w_{t-i}. \quad (43)$$

Define

$$e_t \triangleq u_t^{\text{lin}} - u_t^{\pi(M_q^*)}.$$

Using $w_\tau = 0$ for $\tau \leq 0$ in (41) and (43),

$$e_t = q \sum_{j=H+1}^{t-1} \beta^{j-1} w_{t-j}, \quad (44)$$

where the sum is empty when $t \leq H+1$. Therefore,

$$\begin{aligned} 0 \leq e_t &\leq q \sum_{j=H+1}^{\infty} \beta^{j-1} \\ &= \frac{q\beta^H}{1-\beta}. \end{aligned} \quad (45)$$

State approximation. Let

$$\delta_t \triangleq x_t^{\pi(M_q^*)} - x_t^{\text{lin}}.$$

Subtracting the two state recursions yields

$$\delta_{t+1} = \alpha \delta_t + e_t, \quad \delta_1 = 0. \quad (46)$$

Hence

$$\delta_t = \sum_{s=1}^{t-1} \alpha^{t-1-s} e_s. \quad (47)$$

Combining (45) and (47),

$$\begin{aligned} 0 \leq \delta_t &\leq \frac{q\beta^H}{1-\beta} \sum_{s=1}^{t-1} \alpha^{t-1-s} \\ &\leq \frac{q\beta^H}{(1-\beta)(1-\alpha)}. \end{aligned}$$

Cost approximation. By (5), (45), and (48),

$$\begin{aligned} &\left| c_t(x_t^{\pi(M_q^*)}, u_t^{\pi(M_q^*)}) - c_t(x_t^{\text{lin}}, u_t^{\text{lin}}) \right| \\ &\leq L_x \delta_t + L_u e_t \\ &\leq \frac{L_x q \beta^H}{(1-\beta)(1-\alpha)} + \frac{L_u q \beta^H}{1-\beta} \\ &= \left(L_u + \frac{L_x}{1-\alpha} \right) \frac{q \beta^H}{1-\beta}. \end{aligned} \quad (49)$$

Summing (49) over $t = 1, \dots, T$ and applying the triangle inequality proves (28). $\square$

The remaining lemmas establish the OMD and stability estimates used in the proof of Theorem 1; they rely on the state-decomposition, feasibility, and surrogate-cost lemmas already established in Section IV (Lemmas 1, 2, and 3).

Lemma 5 (Bregman projection onto Δ_H). *Let D_Ψ be the Bregman divergence of the unnormalized negentropy (34). Fix $\widetilde{M} \in \mathbb{R}_{>0}^H$ and set $S \triangleq \sum_{i=1}^H \widetilde{M}^{(i)}$. Then the Bregman projection of $\widetilde{M}$ onto $\Delta_H = \{M \in \mathbb{R}_{\geq 0}^H : \sum_i M^{(i)} \leq 1\}$,*

$$P(\widetilde{M}) \triangleq \arg \min_{M \in \Delta_H} D_\Psi(M \| \widetilde{M}),$$

is given coordinatewise by

$$P(\widetilde{M})^{(i)} = \frac{\widetilde{M}^{(i)}}{\max\{1, S\}}.$$

Proof. Up to the constant $\sum_i \widetilde{M}^{(i)}$, the objective is

$$F(M) \triangleq \sum_{i=1}^H \left(M^{(i)} \log \frac{M^{(i)}}{\widetilde{M}^{(i)}} - M^{(i)} \right).$$

The function F is strictly convex, and, for $M^{(i)} > 0$,

$$\frac{\partial F}{\partial M^{(i)}} = \log \frac{M^{(i)}}{\widetilde{M}^{(i)}}. \quad (50)$$

Its unconstrained minimizer is therefore $M = \widetilde{M}$.

If $S \leq 1$, this minimizer belongs to Δ_H , so $P(\widetilde{M}) = \widetilde{M}$. Suppose $S > 1$. The constraint $\sum_i M^{(i)} \leq 1$ is then active. Moreover, the minimizer has no zero coordinate: if $M^{(k)} = 0$, transferring an arbitrarily small amount of mass from a positive coordinate to coordinate k has directional derivative $-\infty$, because

$$\lim_{z \downarrow 0} \log \frac{z}{\widetilde{M}^{(k)}} = -\infty.$$

Thus the nonnegativity constraints are inactive. With $\lambda \geq 0$ denoting the multiplier of $\sum_i M^{(i)} \leq 1$, stationarity and complementary slackness give

$$\log \frac{M^{(i)}}{\widetilde{M}^{(i)}} + \lambda = 0, \quad \sum_{i=1}^H M^{(i)} = 1. \quad (51)$$

The first relation in (51) implies

$$M^{(i)} = \widetilde{M}^{(i)} e^{-\lambda}.$$

(48) Summing over i and using the second relation in (51),

$$1 = S e^{-\lambda}, \quad e^{-\lambda} = \frac{1}{S}.$$

Hence $\lambda = \log S > 0$ and

$$P(\widetilde{M})^{(i)} = \frac{\widetilde{M}^{(i)}}{S}.$$

Combining the cases $S \leq 1$ and $S > 1$ proves the stated formula. $\square$

Lemma 6 (Closed-form solution of the OMD update). *Fix $M_t \in \Delta_H \cap \mathbb{R}_{>0}^H$, $\eta > 0$, and $g_t \in \mathbb{R}^H$, and let $\widetilde{M}_{t+1}$*

be defined coordinatewise by (36), i.e. $\widetilde{M}_{t+1}^{(i)} = M_t^{(i)} e^{-\eta g_t^{(i)}}$. Then

$$\begin{aligned} & \arg \min_{M \in \Delta_H} \left\{ \eta \sum_{i=1}^H g_t^{(i)} M^{(i)} + D_\Psi(M \| M_t) \right\} \\ &= \arg \min_{M \in \Delta_H} D_\Psi(M \| \widetilde{M}_{t+1}) = P(\widetilde{M}_{t+1}), \end{aligned}$$

where $P(\cdot)$ is the Bregman projection of Lemma 5. Consequently, the OMD update (35) coincides exactly with the closed-form OMD update (36)–(37), and in particular $M_{t+1} \in \Delta_H \cap \mathbb{R}_{>0}^H$. Thus the positive initialization (29) keeps every algorithm iterate strictly positive.

Proof. Fix $M \in \mathbb{R}_{>0}^H$ and a coordinate i , and abbreviate $m = M^{(i)}$, $m_t = M_t^{(i)}$, and $\tilde{m} = \widetilde{M}_{t+1}^{(i)}$. Then

$$\begin{aligned} \eta g_t^{(i)} m + m \log \frac{m}{m_t} &= m \left(\eta g_t^{(i)} + \log \frac{m}{m_t} \right) \\ &= m \log \frac{m}{\tilde{m}}. \end{aligned} \quad (52)$$

Substitution into the OMD objective gives

$$\begin{aligned} & \eta \langle g_t, M \rangle + D_\Psi(M \| M_t) \\ &= D_\Psi(M \| \widetilde{M}_{t+1}) + \sum_{i=1}^H \left(M_t^{(i)} - \widetilde{M}_{t+1}^{(i)} \right). \end{aligned} \quad (53)$$

The final sum is independent of M . Hence the two optimization problems have the same minimizer, and Lemma 5 gives (37). Since $M_t > 0$ coordinatewise, both the multiplicative step and the subsequent positive rescaling preserve strict positivity. $\square$

Lemma 7 (One-step movement of OMD). *For the OMD update with step size $\eta > 0$ and subgradient $g_t \in \partial \ell_t(M_t)$,*

$$\|M_{t+1} - M_t\|_1 \leq \eta \|g_t\|_\infty.$$

In particular, if $\|g_t\|_\infty \leq L$, then $\|M_{t+1} - M_t\|_1 \leq \eta L$.

Proof. Write $\Delta_t \triangleq M_{t+1} - M_t$. First-order optimality of (35), evaluated at the feasible point M_t , gives

$$\begin{aligned} & \langle \eta g_t + \nabla \Psi(M_{t+1}) - \nabla \Psi(M_t), M_t - M_{t+1} \rangle \geq 0, \\ & D_\Psi(M_{t+1} \| M_t) + D_\Psi(M_t \| M_{t+1}) \\ & \leq \eta \langle g_t, M_t - M_{t+1} \rangle \\ & \leq \eta \|g_t\|_\infty \|\Delta_t\|_1. \end{aligned} \quad (54)$$

The second line uses the symmetric Bregman identity.

For every $Z \in \Delta_H \cap \mathbb{R}_{>0}^H$ and $v \in \mathbb{R}^H$,

$$\begin{aligned} v^\top \nabla^2 \Psi(Z) v &= \sum_{i=1}^H \frac{(v^{(i)})^2}{Z^{(i)}} \\ &\geq \frac{\|v\|_1^2}{\sum_{i=1}^H Z^{(i)}} \\ &\geq \|v\|_1^2. \end{aligned} \quad (55)$$

The first inequality is Cauchy–Schwarz. The segment joining M_t and M_{t+1} lies in $\Delta_H \cap \mathbb{R}_{>0}^H$ by Lemma 6; integrating (55) along this segment yields

$$\begin{aligned} D_\Psi(M_{t+1} \| M_t) &\geq \frac{1}{2} \|\Delta_t\|_1^2, \\ D_\Psi(M_t \| M_{t+1}) &\geq \frac{1}{2} \|\Delta_t\|_1^2. \end{aligned}$$

Combining these inequalities with (54),

$$\|\Delta_t\|_1^2 \leq \eta \|g_t\|_\infty \|\Delta_t\|_1.$$

The claim is immediate if $\Delta_t = 0$; otherwise divide by $\|\Delta_t\|_1$. $\square$

Lemma 8 (OMD regret for the surrogate losses). *For every $M \in \Delta_H$, the update (35) satisfies*

$$\sum_{t=1}^T \ell_t(M_t) - \sum_{t=1}^T \ell_t(M) \leq \frac{D_\Psi(M \| M_1)}{\eta} + \frac{\eta}{2} \sum_{t=1}^T \|g_t\|_\infty^2. \quad (56)$$

Proof. Convexity of ℓ_t gives

$$\begin{aligned} \ell_t(M_t) - \ell_t(M) &\leq \langle g_t, M_t - M \rangle \\ &= \langle g_t, M_t - M_{t+1} \rangle \\ &\quad + \langle g_t, M_{t+1} - M \rangle. \end{aligned} \quad (57)$$

Write $h_t \triangleq \nabla \Psi(M_{t+1}) - \nabla \Psi(M_t)$. First-order optimality of (35), evaluated at $M \in \Delta_H$, yields

$$\begin{aligned} \eta \langle g_t, M_{t+1} - M \rangle &\leq \langle h_t, M - M_{t+1} \rangle \\ &= D_\Psi(M \| M_t) - D_\Psi(M \| M_{t+1}) \\ &\quad - D_\Psi(M_{t+1} \| M_t). \end{aligned} \quad (58)$$

The equality is the three-point Bregman identity. Substituting (58) into (57),

$$\begin{aligned} \ell_t(M_t) - \ell_t(M) &\leq \frac{D_\Psi(M \| M_t)}{\eta} - \frac{D_\Psi(M \| M_{t+1})}{\eta} \\ &\quad + \langle g_t, M_t - M_{t+1} \rangle \\ &\quad - \frac{D_\Psi(M_{t+1} \| M_t)}{\eta}. \end{aligned} \quad (59)$$

Let $r_t \triangleq \|M_{t+1} - M_t\|_1$. The strong-convexity calculation (55) and Hölder's inequality give

$$\begin{aligned} & \langle g_t, M_t - M_{t+1} \rangle - \frac{D_\Psi(M_{t+1} \| M_t)}{\eta} \\ & \leq \|g_t\|_\infty r_t - \frac{r_t^2}{2\eta} \\ & \leq \frac{\eta}{2} \|g_t\|_\infty^2. \end{aligned} \quad (60)$$

The last inequality is Young's inequality. Combining (59) and (60), and then summing over t , gives

$$\begin{aligned} \sum_{t=1}^T \ell_t(M_t) - \sum_{t=1}^T \ell_t(M) &\leq \frac{D_\Psi(M \| M_1)}{\eta} \\ &\quad - \frac{D_\Psi(M \| M_{T+1})}{\eta} \\ &\quad + \frac{\eta}{2} \sum_{t=1}^T \|g_t\|_\infty^2. \end{aligned}$$

Since $D_\Psi(M \| M_{T+1}) \geq 0$, this proves (56). $\square$

Lemma 9 (Cumulative movement). *For every $t \in \{2, \dots, T\}$ and every integer i with $1 \leq i \leq t-1$,*

$$\|M_{t-i} - M_t\|_1 \leq i \eta L.$$

Proof. By the triangle inequality,

$$\|M_{t-i} - M_t\|_1 \leq \sum_{j=1}^i \|M_{t-j} - M_{t-j+1}\|_1. \quad (61)$$

Substituting $\|M_{t-j} - M_{t-j+1}\|_1 \leq \eta L$ from Lemma 7 into (61),

$$\begin{aligned} \|M_{t-i} - M_t\|_1 &\leq \sum_{j=1}^i \eta L \\ &= i \eta L. \end{aligned}$$

□

Lemma 10 (Clipping error). *For $M \in \Delta_H$, define*

$$\Gamma_t^M \triangleq x_t^{\pi(M)} - x_t^A. \quad (62)$$

Define the clipping error

$$d_t \triangleq \hat{u}_t - u_t^A. \quad (63)$$

Then, for every $t \in \{1, \dots, T\}$,

$$u_t^A = u_t^{\pi(M_t)} - d_t, \quad 0 \leq d_t \leq \alpha[\Gamma_t^{M_t}]_+. \quad (64)$$

Proof. The clipping rule and the identity $\min\{a, b\} = a - [a - b]_+$ give

$$d_t = [\hat{u}_t - \alpha x_t^A]_+ \geq 0. \quad (65)$$

By (12), $\hat{u}_t = u_t^{\pi(M_t)}$. Fixed-policy feasibility gives

$$\hat{u}_t = u_t^{\pi(M_t)} \leq \alpha x_t^{\pi(M_t)}.$$

Using monotonicity and positive homogeneity of $[\cdot]_+$ in (65),

$$\begin{aligned} d_t &\leq [\alpha x_t^{\pi(M_t)} - \alpha x_t^A]_+ \\ &= \alpha[\Gamma_t^{M_t}]_+. \end{aligned} \quad (66)$$

Finally, (63) and $\hat{u}_t = u_t^{\pi(M_t)}$ give the first identity in (64). □

Lemma 11 (State-gap recursion). *For every fixed $M \in \Delta_H$ and every $t \in \{1, \dots, T\}$,*

$$\Gamma_{t+1}^M = \alpha \Gamma_t^M + \langle M_t - M, \tilde{w}_t \rangle - d_t. \quad (67)$$

Consequently,

$$\Gamma_{t+1}^M \leq \alpha \Gamma_t^M + \langle M_t - M, \tilde{w}_t \rangle, \quad (68)$$

$$\Gamma_{t+1}^M \geq \alpha \Gamma_t^M + \langle M_t - M, \tilde{w}_t \rangle - \alpha[\Gamma_t^{M_t}]_+. \quad (69)$$

Proof. The two trajectories use the same arrival w_t . Hence

$$\begin{aligned} \Gamma_{t+1}^M &= \alpha \Gamma_t^M - u_t^{\pi(M)} + u_t^A \\ &= \alpha \Gamma_t^M - u_t^{\pi(M)} + u_t^{\pi(M_t)} - d_t \\ &= \alpha \Gamma_t^M + \langle M_t - M, \tilde{w}_t \rangle - d_t. \end{aligned} \quad (70)$$

This proves (67). The two bounds follow from $0 \leq d_t \leq \alpha[\Gamma_t^{M_t}]_+$ in Lemma 10. □

Lemma 12 (State deviation bound). *For every $t \in \{1, \dots, T\}$, with Γ_t^M as in (62),*

$$-\frac{\alpha \eta L}{(1-\alpha)^3} \leq \Gamma_t^{M_t} \leq \frac{\alpha \eta L}{(1-\alpha)^2}.$$

Proof. Fix $t \in \{1, \dots, T\}$. For $s < t$, define

$$\xi_{s,t} \triangleq \langle M_s - M_t, \tilde{w}_s \rangle.$$

Since $\tilde{w}_s^{(i)} = \alpha^i w_{s-i}$,

$$\|\tilde{w}_s\|_\infty \leq \alpha.$$

Lemma 9 and Hölder's inequality therefore give

$$\begin{aligned} |\xi_{s,t}| &\leq \|M_s - M_t\|_1 \|\tilde{w}_s\|_\infty \\ &\leq \alpha \eta L(t-s). \end{aligned} \quad (71)$$

For the fixed comparator M_t , unrolling (67) from $\Gamma_1^{M_t} = 0$ gives

$$\Gamma_t^{M_t} = \sum_{s=1}^{t-1} \alpha^{t-1-s} (\xi_{s,t} - d_s). \quad (72)$$

Since $d_s \geq 0$, (71) implies

$$\begin{aligned} \Gamma_t^{M_t} &\leq \sum_{s=1}^{t-1} \alpha^{t-1-s} |\xi_{s,t}| \\ &\leq \alpha \eta L \sum_{s=1}^{t-1} \alpha^{t-1-s} (t-s) \\ &= \alpha \eta L \sum_{j=1}^{t-1} j \alpha^{j-1} \\ &\leq \frac{\alpha \eta L}{(1-\alpha)^2}. \end{aligned} \quad (73)$$

The last line uses $\sum_{j=1}^{\infty} j \alpha^{j-1} = (1-\alpha)^{-2}$.

Applying (64) and (73) at time s ,

$$\begin{aligned} 0 \leq d_s &\leq \alpha[\Gamma_s^{M_s}]_+ \\ &\leq \frac{\alpha^2 \eta L}{(1-\alpha)^2}. \end{aligned} \quad (74)$$

This use of (73) is noncircular: its derivation used only $d_s \geq 0$. From (72), (71), and (74),

$$\begin{aligned} \Gamma_t^{M_t} &\geq -\sum_{s=1}^{t-1} \alpha^{t-1-s} |\xi_{s,t}| - \sum_{s=1}^{t-1} \alpha^{t-1-s} d_s \\ &\geq -\frac{\alpha \eta L}{(1-\alpha)^2} - \frac{\alpha^2 \eta L}{(1-\alpha)^2} \sum_{s=1}^{t-1} \alpha^{t-1-s} \\ &\geq -\frac{\alpha \eta L}{(1-\alpha)^2} - \frac{\alpha^2 \eta L}{(1-\alpha)^3} \\ &= -\frac{\alpha \eta L}{(1-\alpha)^3}. \end{aligned} \quad (75)$$

Together, (73) and (75) prove the claim. □

Lemma 13 (Action gap induced by clipping). *For all $t \in \{1, \dots, T\}$,*

$$0 \leq u_t^{\pi(M_t)} - u_t^A \leq \frac{\alpha^2 \eta L}{(1-\alpha)^2}.$$

Proof. By (64),

$$u_t^{\pi(M_t)} - u_t^A = d_t \geq 0.$$

The same bound and Lemma 12 give

$$\begin{aligned} d_t &\leq \alpha[\Gamma_t^{M_t}]_+ \\ &\leq \frac{\alpha^2 \eta L}{(1-\alpha)^2}. \end{aligned}$$

□

Lemma 14 (Stability term). *For all horizons T ,*

$$\sum_{t=1}^T \left(c_t(x_t^A, u_t^A) - c_t(x_t^{\pi(M_t)}, u_t^{\pi(M_t)}) \right) \leq \frac{\alpha T \eta L^2}{(1-\alpha)^2}.$$

Proof. Both state-action pairs lie in $\mathcal{D}_\alpha$. Hence (5), Lemma 12, and Lemma 13 give

$$\begin{aligned} &c_t(x_t^A, u_t^A) - c_t(x_t^{\pi(M_t)}, u_t^{\pi(M_t)}) \\ &\leq L_x |\Gamma_t^{M_t}| + L_u |u_t^A - u_t^{\pi(M_t)}| \\ &\leq \frac{\alpha \eta L L_x}{(1-\alpha)^3} + \frac{\alpha^2 \eta L L_u}{(1-\alpha)^2} \\ &= \frac{\alpha \eta L}{(1-\alpha)^2} \left(\frac{L_x}{1-\alpha} + \alpha L_u \right) \\ &= \frac{\alpha \eta L^2}{(1-\alpha)^2}. \end{aligned} \quad (76)$$

Summing (76) over $t = 1, \dots, T$ proves the claim. □

Proof of Theorem 1. Define the stability term

$$S_T \triangleq \sum_{t=1}^T (c_t(x_t^A, u_t^A) - \ell_t(M_t)).$$

Using $\ell_t(M) = c_t(x_t^{\pi(M)}, u_t^{\pi(M)})$, the regret decomposes as

$$\text{Reg}_{\text{SDAC}}(T) = S_T + \sum_{t=1}^T \ell_t(M_t) - \min_{M \in \Delta_H} \sum_{t=1}^T \ell_t(M). \quad (77)$$

We first bound the initial Bregman divergence. Fix $M \in \Delta_H$ and set $W \triangleq \sum_{i=1}^H M^{(i)} \in [0, 1]$. Since $M_1^{(i)} = 1/(H+1)$,

$$\begin{aligned} D_\Psi(M \| M_1) &= \sum_{i=1}^H M^{(i)} \log((H+1)M^{(i)}) \\ &\quad - W + \frac{H}{H+1} \\ &= \sum_{i=1}^H M^{(i)} \log M^{(i)} + W \log(H+1) \\ &\quad - W + \frac{H}{H+1}. \end{aligned} \quad (78)$$

Define the affine function

$$f_H(W) \triangleq W \log(H+1) - W + \frac{H}{H+1}.$$

Because $z \log z \leq 0$ for $z \in [0, 1]$,

$$D_\Psi(M \| M_1) \leq f_H(W). \quad (79)$$

The function f_H is affine on $[0, 1]$, and

$$\begin{aligned} f_H(0) &= \frac{H}{H+1} \leq \log(H+1), \\ f_H(1) &= \log(H+1) - \frac{1}{H+1} \leq \log(H+1). \end{aligned}$$

Here the first inequality follows from $\log z \geq 1 - z^{-1}$ for $z \geq 1$. Therefore,

$$D_\Psi(M \| M_1) \leq \log(H+1) \quad \text{for every } M \in \Delta_H. \quad (80)$$

Lemma 8, $\|g_t\|_\infty \leq L$, and (80) imply, for every $M \in \Delta_H$,

$$\begin{aligned} \sum_{t=1}^T (\ell_t(M_t) - \ell_t(M)) &\leq \frac{D_\Psi(M \| M_1)}{\eta} + \frac{\eta}{2} \sum_{t=1}^T \|g_t\|_\infty^2 \\ &\leq \frac{\log(H+1)}{\eta} + \frac{\eta L^2 T}{2}. \end{aligned} \quad (81)$$

Since (81) holds uniformly over M ,

$$\sum_{t=1}^T \ell_t(M_t) - \min_{M \in \Delta_H} \sum_{t=1}^T \ell_t(M) \leq \frac{\log(H+1)}{\eta} + \frac{\eta L^2 T}{2}. \quad (82)$$

Lemma 14 gives

$$S_T \leq \frac{\alpha \eta L^2 T}{(1-\alpha)^2}. \quad (83)$$

Substituting (82) and (83) into (77),

$$\text{Reg}_{\text{SDAC}}(T) \leq \frac{\log(H+1)}{\eta} + \eta L^2 T \left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right). \quad (84)$$

Set

$$A \triangleq \log(H+1), \quad B \triangleq L^2 T \left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right).$$

Both constants are positive. The right-hand side of (84) is $A/\eta + B\eta$, whose derivative is

$$-\frac{A}{\eta^2} + B.$$

Its unique minimizer is therefore

$$\eta^* = \sqrt{\frac{A}{B}} = \sqrt{\frac{\log(H+1)}{L^2 T \left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right)}}. \quad (85)$$

At this value,

$$\frac{A}{\eta^*} = \sqrt{AB}, \quad B\eta^* = \sqrt{AB}.$$

Thus (84) becomes

$$\text{Reg}_{\text{SDAC}}(T) \leq 2L \sqrt{\left(\frac{1}{2} + \frac{\alpha}{(1-\alpha)^2} \right) T \log(H+1)}. \quad (86)$$

□

REFERENCES

- [1] N. Agarwal, B. Bullins, E. Hazan, S. Kakade, and K. Singh, "Online control with adversarial disturbances," in *International Conference on Machine Learning*. PMLR, 2019, pp. 111–119.
- [2] N. Agarwal, E. Hazan, and K. Singh, "Logarithmic regret for online control," *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [3] M. Simchowitz, K. Singh, and E. Hazan, "Improper learning for non-stochastic control," in *Conference on Learning Theory*. PMLR, 2020, pp. 3320–3436.
- [4] E. Hazan and K. Singh, "Introduction to Online Control," Apr. 2026, arXiv:2211.09619, revised Apr. 2026. [Online]. Available: <https://arxiv.org/abs/2211.09619>

- [5] S. Lale, K. Azizzadenesheli, B. Hassibi, and A. Anandkumar, "Logarithmic regret bound in partially observable linear dynamical systems," *Advances in Neural Information Processing Systems*, vol. 33, pp. 20 876–20 888, 2020.
- [6] E. Hazan, S. M. Kakade, K. Singh, and A. Sidford, "Nonstochastic control with bandit feedback," *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, pp. 16 774–16 785, 2020.
- [7] Y. Li, S. Das, and N. Li, "Online optimal control with affine constraints," *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, no. 10, pp. 8527–8537, 2021. [Online]. Available: <https://ojs.aaai.org/index.php/AAAI/article/view/17035>
- [8] X. Liu, Z. Yang, and L. Ying, "Online nonstochastic control with adversarial and static constraints," in *Proceedings of the 40th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 22 277–22 288. [Online]. Available: <https://proceedings.mlr.press/v202/liu23at.html>
- [9] Q. Lin, J. Su, and M. Chen, "Optimal algorithms for online Age-of-Information optimization in energy harvesting systems," *IEEE Transactions on Networking*, vol. 33, no. 6, pp. 3146–3161, 2025.
- [10] H. Yu and M. J. Neely, "Learning-aided optimization for energy-harvesting devices with outdated state information," *IEEE/ACM Transactions on Networking*, vol. 27, no. 4, pp. 1501–1514, 2019.
- [11] K. Asgari and M. J. Neely, "Bregman-style online convex optimization with energy harvesting constraints," *Proc. ACM Meas. Anal. Comput. Syst.*, vol. 4, no. 3, nov 2020. [Online]. Available: <https://doi.org/10.1145/3428337>
- [12] L. Tassiulas and A. Ephremides, "Stability properties of constrained queueing systems and scheduling policies for maximum throughput in multihop radio networks," in *29th IEEE Conference on Decision and Control*. IEEE, 1990, pp. 2130–2132.
- [13] —, "Dynamic server allocation to parallel queues with randomly varying connectivity," *IEEE Transactions on Information Theory*, vol. 39, no. 2, pp. 466–478, 1993.
- [14] N. McKeown, A. Mekkittikul, V. Anantharam, and J. Walrand, "Achieving 100% throughput in an input-queued switch," *IEEE Transactions on Communications*, vol. 47, no. 8, pp. 1260–1267, 1999.
- [15] M. J. Neely, *Stochastic Network Optimization with Application to Communication and Queueing Systems*. Morgan & Claypool Publishers, 2010.
- [16] M. A. Khojastepour and A. Sabharwal, "Delay-constrained scheduling: Power efficiency, filter design, and bounds," in *IEEE INFOCOM 2004*, vol. 3. IEEE, 2004, pp. 1938–1949.
- [17] L. Yang, M. H. Hajiesmaili, R. Sitaraman, A. Wierman, E. Mallada, and W. S. Wong, "Online linear optimization with inventory management constraints," *Proceedings of the ACM on Measurement and Analysis of Computing Systems*, vol. 4, no. 1, 2020.
- [18] M. Hihat and A. Fermanian, "Online policy selection for inventory problems," *arXiv preprint arXiv:2411.19269*, 2024. [Online]. Available: <https://arxiv.org/abs/2411.19269>
- [19] P. A. P. Moran, "A probability theory of dams and storage systems," *Australian Journal of Applied Science*, vol. 5, no. 2, pp. 116–124, 1954.
- [20] N. U. Prabhu, *Stochastic Storage Processes: Queues, Insurance Risk, Dams, and Data Communication*. New York, NY: Springer New York, 1998. [Online]. Available: <https://doi.org/10.1007/978-1-4612-1742-8>
- [21] E. Hazan, "Introduction to online convex optimization," *Foundations and Trends in Optimization*, vol. 2, no. 3–4, pp. 157–325, 2016.